

Pan & Tilt

Regulering og Kontrol af Robotssystemer - Group 5

Education and Semester:

Robotic engineering - 4. semester

Date of Submission:

25-05-2016

Supervisor:

Sarah Madeleine Garney

Group members:

Mathias Gregersen, Anton M.M. Schroll,

Daniel Haraldson, René T Haagensen,

Anders Ellinge, Bosse Hougesen

& Steen-Bjørn C.R. Vinther

Technical Faculty

University of Southern Denmark

SpotLight

Preface

This report is written by seven 4th semester robotic engineering students from the University of Southern Denmark in spring 2016. The report is based on the 4th semester's project. It is aimed at people with basic knowledge in the following disciplines: Digital programmable electronics, Embedded Programming, Control Engineering and Computer systems. Furthermore, in order to understand the contents of this report, it is considered a prerequisite that the project introduction has been read. This can be found in appendix B. This report describes the theory and analysis involved in designing the controller used to control the Pan & Tilt system, as well as the motor driver and the application used to give commands to the system. The report is written as a conclusion to the project and contains documentation of the process and the experiences gained. It should therefore be possible to recreate the project using this report.

Credentials of group members

Name - SDU Mail - Exam Number

Bosse Hougesen - bohoul3@student.sdu.dk - 334023

Daniel Haraldson - dahar14@student.sdu.dk - 66651375

Anders Ellinge - aelli14@student.sdu.dk - 65082066

Anton Schroll - aschr13@student.sdu.dk - 334691

Steen-Bjørn Vinther - stvin13@student.sdu.dk - 336007

Mathias Gregersen - magre14@student.sdu.dk - 409938

René Haagensen - rehaa14@student.sdu.dk - 412515

1 Abstract

The purpose of this paper is to show how to make a controller and how to use it to control a spotlight consisting of two connected frames, two motors and one light. Precise motor control is an essential requirement in most applications involving any kind of actuators that need to behave in a predictable and controlled way. A spotlight capable of relative precise movement is important when performing on a stage. In order to make a controller, several measurements of the systems response to different voltages have been analysed. Using Matlab and control theory, the values of the controller have been estimated. The results show; with the right controller, the system is easily controllable and will behave in a predictable way.

Table of content

1	Abstract	3
2	Introduction	7
2.1	Object Balancing	7
2.2	Spotlight Control	7
2.3	Triangulation	7
2.4	Basis of Evaluation	7
3	Problem formulation	8
3.1	Definition of goals	8
3.2	Requirements specification	8
3.2.1	Primary requirements	8
3.2.2	Secondary requirements	9
4	Overview of the System	9
5	SPI communication	11
5.1	SPI communication protocol	11
5.2	Addressing	12
6	The Plant	13
6.1	DC Motor Modelling	13
6.1.1	DC Motor Tests	13
6.1.2	Discussion	15
6.1.3	SI-units for DC Gain	17
6.2	P&T System Modelling	17
6.2.1	Inertia in the P&T system	17
6.2.2	Tests of the P&T motors	18
6.3	Overview	19
7	Control Theory	20
7.1	Introduction	20
7.2	Controller selection	21
7.2.1	Proportional Controller	21
7.2.2	Proportional-Integral Controller	22
7.2.3	Proportional-Derivative Controller	23
7.2.4	Proportional-Integral-Derivative Controller	23
7.3	K constants	24
7.3.1	Manual tuning	25
7.3.2	Root Locus tuning method	25
7.3.3	Controller constants for Pan and Tilt systems	25
7.4	Simulation	26
7.5	Implementation - digital	30

7.5.1	Direct implementation	30
7.5.2	Filter approach	30
7.5.3	Discrete filter for Pan-system [10kHz]	30
7.5.4	Discrete filter for Tilt-system [10kHz]	31
7.6	Overview	31
8	The Microcontroller	31
8.1	Choice of Operating System	32
8.2	Program Structure	32
8.2.1	Defining tasks	32
8.2.2	Connecting the tasks	33
8.3	The Application Programming Interface	33
8.4	Event Queues and Shared State Memory	35
8.5	Deadlock Protection	35
8.6	SPI-Master	36
8.6.1	Performance Test	36
8.7	The Application	37
8.8	Application Conclusion	37
9	Driver in the FPGA	38
9.1	Overall analysis	38
10	SPI slave block	41
10.1	Description	41
10.2	Implementation	41
10.3	Tests	43
10.4	Discussion	43
10.5	Conclusion	43
11	Controller block	43
11.1	Description	43
11.2	Mathematical techniques	44
11.2.1	Discrete Integration	44
11.2.2	Shortest path	44
11.3	Implementation	45
11.3.1	Comparator	45
11.3.2	PI controller	46
11.4	Tests	47
11.4.1	Comparator	47
11.4.2	PI controller	47
11.5	Discussion	47
11.6	Conclusion	47
12	Sensor block	48

12.1 Description	48
12.2 Theory	48
12.3 Implementation	49
12.4 Test	49
12.5 Discussion	50
12.6 Conclusion	50
13 Motor driver block	50
13.1 Description	50
13.2 Theory	50
13.2.1 PWM generation	51
13.2.2 H-bridge	51
13.3 Implementation	51
13.3.1 PWM part	51
13.3.2 Direction part	52
13.4 Tests	52
13.5 Discussion	52
13.6 Conclusion	52
14 Driver conclusion	52
15 Conclusion	53
16 Appendix	54
A Kernel Instruction List	54
B Content on CD	54

2 Introduction

In order to set limits and generate criteria for the realization of the project, the group decided to give the project an application, rather than just designing a control system for the P&T setup. The following ideas were considered and graded.

2.1 Object Balancing

An object is placed on a plate attached to the P&T system. The core of this project is the balancing of the object, making sure it remains still on the plate. It should also be possible to move the plate to make the object move to a specific location on the surface. Expansion: Furthermore, a labyrinth could be designed for the P&T system to navigate a ball through.

2.2 Spotlight Control

A spotlight is attached to the system, making it adjustable without physically moving the system. The P&T system is used to point the spotlight to the desired position. This can be done by either giving the system coordinates or by creating an application that moves the spotlight through predefined patterns.

2.3 Triangulation

This would involve the triangulation of a signal, using the P&T movement to point something in the direction of it. An example use for this could be a camera attached to the P&T system, recording a person carrying a signaling device.

2.4 Basis of Evaluation

Choice of the project is based on the following criteria:

1. Difficulty
 - (a) Is the project reasonable in respect to making early progress?
 - (b) Is the project sufficiently complex?
2. Interest - Does the project seem interesting?
3. Curriculum relevance - To what degree does the project relate to the courses of the semester?
4. Practical relevance

The value below the criteria titles is the weighting. Larger values indicate higher relevance.

Through this evaluation, the spotlight control project was decided upon. From this point, every decision was made with respect to this selection.

Project name	Difficulty [4]	Interest [2]	Curr. rel. [4]	Practical rel. [1]	Sum
Object balance	4	4,5	5	3	48
Spotlight control	5	5	5	4	54
Triangulation	3	3,5	3	5	36

Figure 1: Table evaluating the merit of different project ideas

3 Problem formulation

When performing on a stage, proper handling of light is important. This can be achieved through controllable spotlights. They have to be able to rotate horizontally and vertically, thereby targeting specific points on the stage and do so autonomously.

The Pan & Tilt System (P & T System) can be made to meet these requirements.

3.1 Definition of goals

The main goals of the product is, to be able to point the cone of light from the spotlight to a specific point on a predefined surface, or follow a human walking around on a surface. The user should have the option to specify the limits of the spotlights working space. The positions or paths should be set by the user and the movement of the cone should be pleasant to watch.

3.2 Requirements specification

The requirements will be divided into primary and secondary requirements, so as to identify the most important objectives.

3.2.1 Primary requirements

Hardware requirements

- The controller has to be implemented on the FPGA(Field-programmable gate array).
- The communication between the microprocessor and the FPGA has to be done with SPI(Serial Peripheral Interface Bus).
- The FPGA has to control the PWM(Pulse-width modulation) signal for the motors.

Driver requirements

- The system cannot have an overshoot of more than 10%.
- An increment in the hall sensor counter must have a settling time of at least 50 mS.

- The spotlight should have a steady state close to zero

Application requirements

- The application must have support for the different spotlight height configurations.
- The application must support different stage dimensions, to a maximum of 14 m x 14 m.

3.2.2 Secondary requirements

These are the requirements that would be nice improvements to the project, but are not strictly necessary for the successful completion of the project.

- The movement between two points should be in a smooth motion.
- The P&T system is to recalibrate the position of the Pan and the tilt during startup.
- The spotlight can be controlled by an application.
- The spotlight can be controlled by a joystick.
- At a scene of dimension 5 m x 5 m, from a height of 5 meters, the system must achieve a velocity of at least 2.5m/s^1 relative to the scene, which is the maximum velocity the average person would walk. ¹

4 Overview of the System

In the process of brainstorming, it was a necessity to specify the responsibility of the system, and to create clear interfaces between the modules and components. In order to create a better overview of the system, a block diagram was made (figure 2) to illustrate all the parts and peripherals of the system.

The block diagram shows 3 main blocks; the application, the P&T driver and the hardware. The application is the user's way of controlling the system, through the use of either the joystick or the console commands. This block only contains the necessary code to receive input, and to communicate with the P&T API, which is done through inter-task communication. The P&T driver is divided into 2 sub blocks: The microcontroller and the FPGA. In this system, a clear split between software and hardware is made, with the microcontroller containing all the software-related code, and the FPGA containing all the hardware-related code. The microcontroller block contains simple C code and memory, and passes necessary information to the FPGA via SPI. This block generally contains code which is non-critical to performance. The FPGA block contains the controller for the P&T system, and sends PWM and direction to the P&T hardware. The advantages of placing the controller in the FPGA, is getting quicker response from the hardware, as well as lower latency compared to placing the controller in the microcontroller. The hardware block contains the physical hardware, the P&T system, and the feedback from the tachometers. This block receives input and returns feedback to the FPGA.

¹Wikipedia, Preferred walking speed, https://en.wikipedia.org/wiki/Preferred_walking_speed, 23/05/16

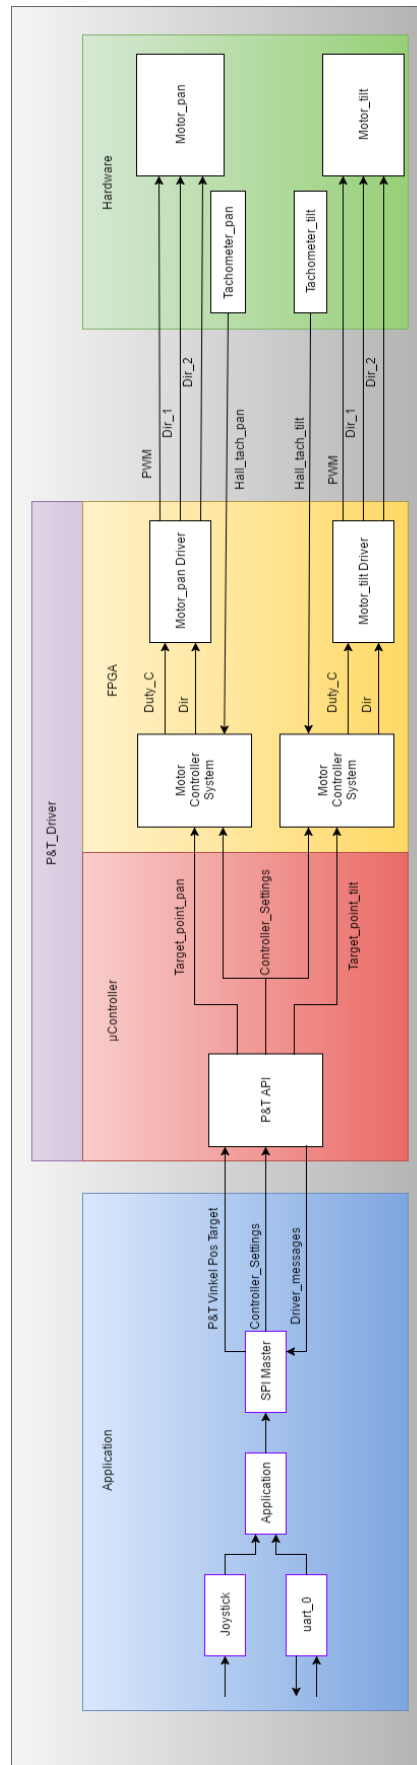


Figure 2: The overall system block diagram

5 SPI communication

This section will explain how the connection between the FPGA and the microcontroller is made via an SPI communication protocol. This section can be divided into two parts with the first part concerning the general SPI communication and the second part concerning addressing of the data sendt. Note that the second part, in principle, has nothing to do with the actual SPI communication, however since it is important for both the FPGA microcontroller it makes sense to talk about it before going into detail with the FPGA and microcontroller.

5.1 SPI communication protocol

The SPI communication protocol is used to send data, one bit at a time. The SPI communication protocol runs a master slave relationship which means that one of the parts involved is the master part and fully controls the connection and the other parts simply does what the master part tells them to do. An example of this can be seen on figure 3.

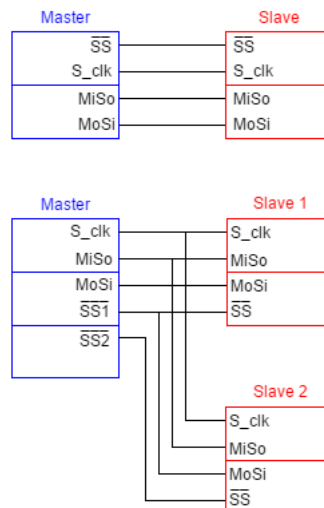


Figure 3: Diagram of how the master and slave is connected

In order to send the data four wires, called Slave Select (SS), Master out Slave in (MOSI), Master in Slave out (MISO) and Serial Clock (S.CLK), are used. The MISO and MOSI wires is where the data bits are through and SS and S.CLK are wires that controls the transfer, note that SS and S.CLK is controlled only by the master. When the SS is active a data transfer is initiated between the master and a slave, the master is telling the slave that they are now communicating. When the SS is deactivated, the transfer is over and all of the data is sent. Note that the SS is active low, which means that it is active when the SS signal is low and deactivated when the signal is high.

The S_CLK is used to signal a read of the MOSI and MISO line. This is very important since this is where the synchronization of the data transfer is made. Without this synchronization the data transfer would be impossible. An example of a SPI data transfer can be seen on figure 4.

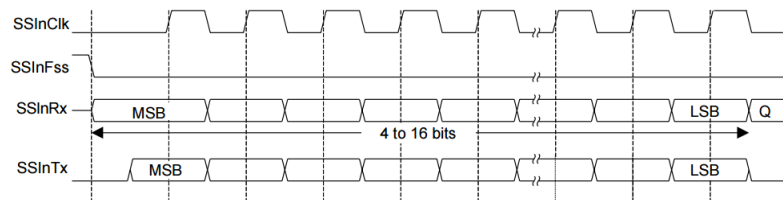


Figure 4: Illustration of how the package is constructed

5.2 Addressing

Instead of running several SPI slaves in the communication to send the data to the right part of the FPGA, giving the data an address instead was chosen. This makes the amount of data sent longer, however since the controller is a part of the FPGA the speed which information is sent does not matter. So since the addressing was found to be the easier solution to make it was chosen. Addressing works by appending extra bits onto the data, much like a header. An example of this can be seen on figure 5. The address bits then define what kind of data is being sent and/or where it has to go. In the case of the SPI slave in the FPGA, the address symbolizes where the data has to go, since what the data symbolizes is an information not needed in the FPGA. In the case of the SPI master in the microcontroller, the address symbolizes what kind of data was received.

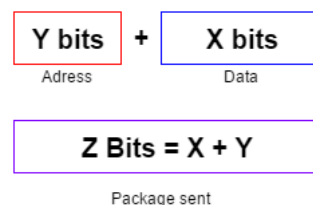


Figure 5: Illustration of how the package is constructed

When receiving a packet via the SPI communication, the receiver separates the address from the data and then uses the data according to what the address tells about the data. An example of this can be seen on figure 6.

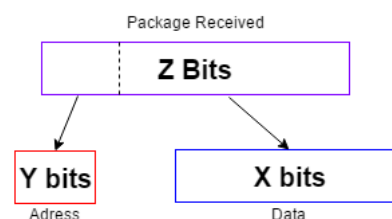


Figure 6: Illustration of how the package is deconstructed

6 The Plant

In order to properly understand and control the P&T system, mathematical models for the pan and tilt systems (the plant) had to be developed. This section will discuss the process of initially testing for relevant DC motor data, the derivation of a mathematical model, then comparing this model to the test data for validation of its accuracy. Finally, the same method will be used to develop models for the behaviour of the pan and tilt systems as a whole.

6.1 DC Motor Modelling

6.1.1 DC Motor Tests

The purpose of these tests are to analyse and generate a blackbox model for the DC motor, based on its step response.

The motor was set up with a switch connected to a power supply set at the desired voltage to test. This would create a physical step response from the motor when the switch was toggled. The two encoders output on the EMG30 are open collectors, so the outputs were connected to pull-up resistance to produce a desired signal. The pull-up resistors for these tests were picked to be around 300 ohm.

An oscilloscope was set to measure the two encoder outputs on their own individual channels. The sampling starts on an edge trigger with a sampling rate of 2 MHz and stops the sampling after 0.5 seconds. These 2 set of samples would together be the data defining an step response.

Following step responses were sampled (Figure 7), and the information extracted (Figure 8) from the data was calculated in matlab.




Figure 7: DC motor step responses from 1 to 12V. The bottom red line is 1V, and the upper yellow line is 12V

No Load\Volt	Measured RPM	Rated RPM	Offset	DC gain	Time constant Tau	Tachs/s
1	10,94	18	3,269 %	10,9	0,0185	65,64
2	28,92	36	3,278 %	14,5	0,0111	173,52
3	46,9	54	3,287 %	15,6	0,005	281,4
4	66,5	72	2,546 %	16,6	0,0085	399
5	83,5	90	3,009 %	16,7	0,0077	501
6	103,3	108	2,176 %	17,2	0,0053	619,8
7	123,4	126	1,204 %	17,6	0,0096	740,4
8	141,4	144	1,204 %	17,7	0,0068	848,4
9	160,2	162	0,833 %	17,8	0,0087	961,2
10	179,4	180	0,278 %	17,9	0,0075	1076,4
11	198,1	198	-0,046 %	18,0	0,0081	1188,6
12	216,4	216	-0,185 %	18,0	0,0108	1298,4
Avg			1,738 %	16,6	0,0090	

Figure 8: Table summarizing the information from the DC motor tests.

As noted, the step responses of the motor appear to be of a first-order linear-time-invariant (LTI) system. This means the system can be estimated as:

$$DCgain \frac{1}{\tau s + 1} \quad (1)$$

The DC gain is the steady state value divided by the steady state (constant) value of the input step.

$$\frac{216}{12V} = 18 \frac{rpm}{V} \quad (2)$$

This indicates that RPM is directly proportional to a voltage. The expected proportionality (in blue) is plotted below together with the actual measured RPM (in green) in figure 9.

As seen, the actual RPM of the motor is not directly proportional to a voltage, but the difference is at max a 3.287% offset from the expected. Since this is relatively low, the system is estimated to be proportional. Though at low voltages the relative DC gain at that voltage starts to differ significantly from the estimated values. Therefore the DC gain is estimated as the mean value of the actual DC gains measured for the voltages 4V to 12V, since the DC gain in this range is relatively linear. The DC gain is: $\approx 17.6 \frac{rpm}{V}$

The time constant tau is the time when the step response reaches $1 - 1/e \approx 63.2\%$ of its final value. This time constant is expected to be constant, but as seen in figure 10, at low voltages there seems to be some uncertainties regarding the time constant.

This behavior can be explained by variations in the actual experiments, since the encoders are not uniformly placed in the motor. The start position of the encoders can have some impact regarding the experimental results. Other than that, it would be expected to see some differences from the lower voltages compared to the higher, since the DC gain of the system indicated some nonlinearity. The time constant, Tau, is then estimated to be the mean value from 1 to 12 volts, $\tau \approx 0.009seconds$.

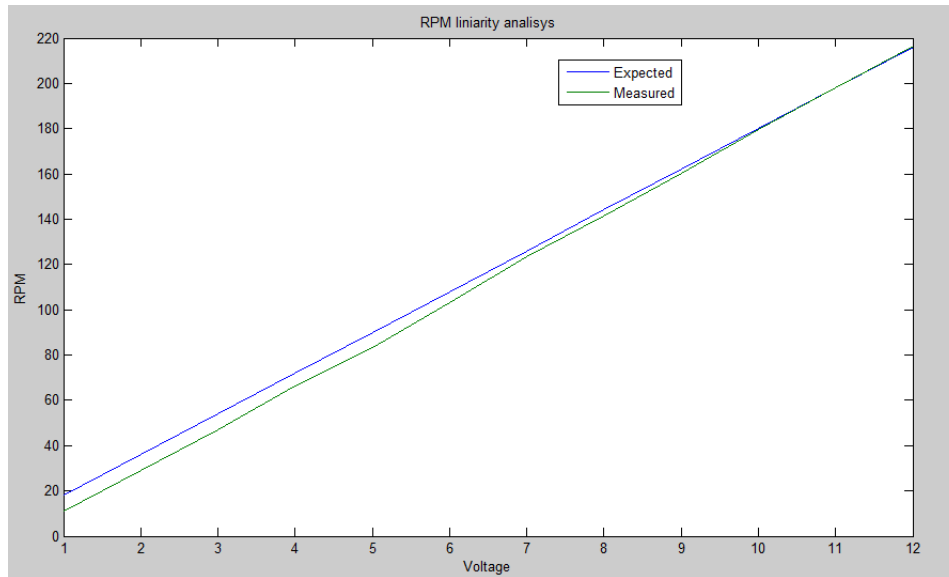


Figure 9: Comparing measured RPM values to expected ones




Figure 10: Time constant for the DC Motor at different voltages

The transfer function for the DC motor is therefore now defined as follows:

$$G(s) = \frac{17.6}{0.009s + 1} \quad (3)$$

As seen on Figure 11 the first-order approximation $G(s)$ of the motor is relatively accurate.

6.1.2 Discussion

Having modelled our DC-motor as a first order LTI system, it is relevant to discuss the reliability and accuracy of this model. In reality, the transfer function for the motor includes both an external, mechanical part and an

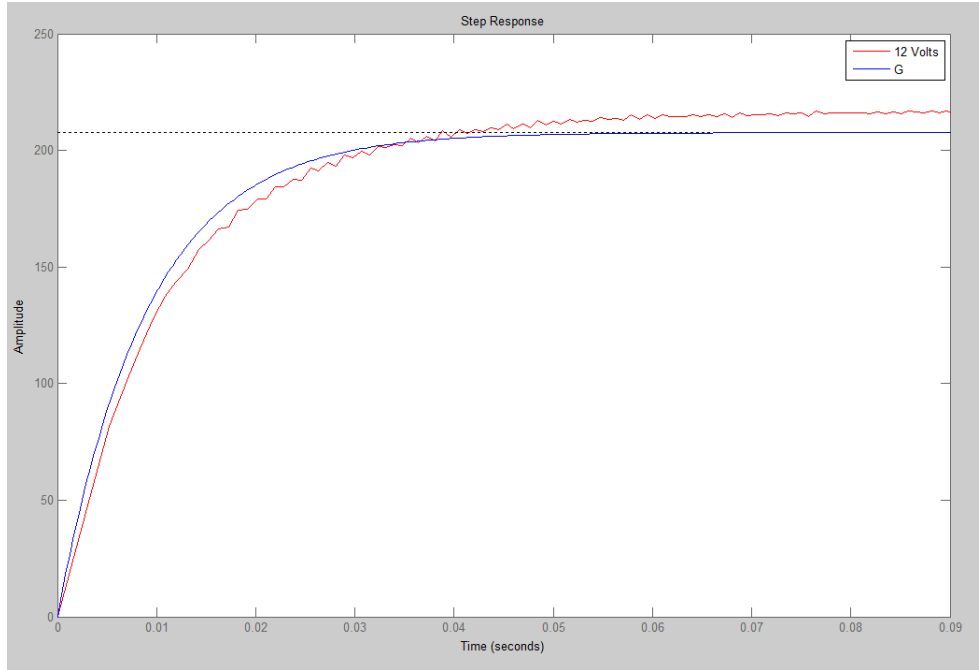


Figure 11: Comparison of the first order approximation step response and the actual step response from research.

internal, electrical part (the armature).²

$$\frac{\theta(s)}{V_f(s)} = G(s) = \frac{K_m/(bR_f)}{s(\tau_f s + 1)(\tau_L s + 1)} \quad (4)$$

This model is actually a second order system. However the time constant τ_f , denoting the response time of the internal electronics of the motor, is usually negligible compared to the mechanical dynamics of the motor³. This makes sense both intuitively and mathematically, as the relatively faster armature with an associated pole further away from origin will have less of an impact on the overall system response, and can thus be omitted for the sake of simplicity. Typical values for these time constants show this as well.

The first and second order transfer functions generated using these time constants are close to identical, as can be seen in figure 12. The motor constants used:

Field time constant: $\tau_f = 1ms$

Rotor time constant⁴: $\tau = 100ms$

Omitting the armature from the derived transfer function gives us the following, which essentially equals our first order estimation. This indicates that a first order approximation of the motor can be very accurate, while simultaneously simplifying the controller design process.

$$\frac{K_m/(R_a b + K_b K_m)}{s(\tau_1 s + 1)} \quad (5)$$

²Model and transfer function from Control Systems: Twelfth Edition p. 72

³Modern Control Systems: Twelfth Edition p72

⁴Modern Control Systems: Twelfth edition p77.

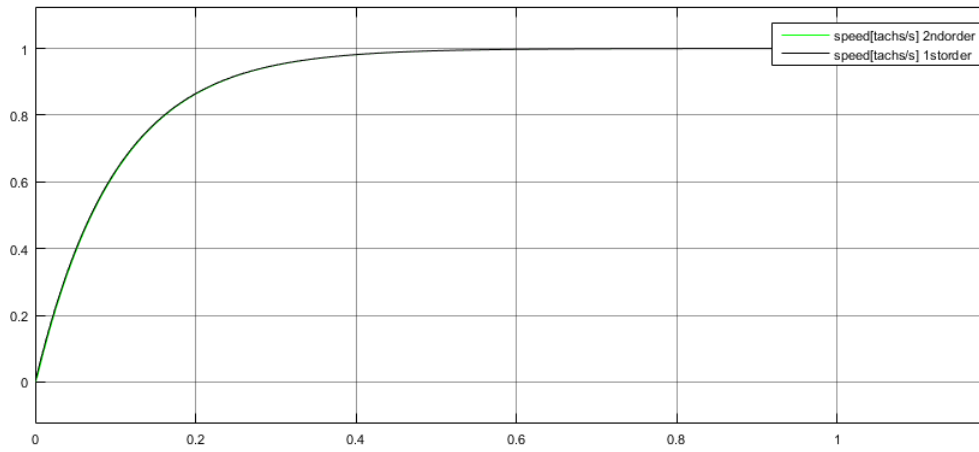


Figure 12: Comparing 1st and 2nd order transfer function step responses. x-axis: sec, y-axis: normalized DC gain

6.1.3 SI-units for DC Gain

Looking at the model in the larger context of our control system, it is evident that using RPM/V as the SI-unit for DC-gain is less than optimal. The feedback from the Hall sensors is in the form of tachs, and would thus require a conversion in order to be interpreted as motor rotations. In addition to that, the controlled variable of our system is the position of the motor in terms of tachs. The transfer function can easily be modified to reflect this through an integration, effectively making our transfer function a second order system.

Speed:

$$G(s) = \frac{105.6}{0.009s + 1} \frac{[tachs/sec]}{[V]} \quad (6)$$

Position:

$$G(s) = \frac{105.6}{0.009s^2 + s} \frac{[tachs/sec]}{[V]} \quad (7)$$

The behavior of these transfer functions given a step input can be seen in figure 13. This behavior looks very promising. The speed transfer function step response shows an acceleration, until the speed reaches the $DCGain * 12V$. The position transfer function also behaves as expected, as the area below the speed transfer function.

6.2 P&T System Modelling

Having modelled the DC motor, it is now time to move on to the modelling of the entire plant. The same methods will be used, except now we are testing and considering the P&T loads and their influence on the system.

6.2.1 Inertia in the P&T system

Because of inertia present in the P&T system parts, it can be expected that the constants of the individual motors will look different than those describing the motor without load. Also, because transfer function of the DC motor without load is approximated to be linear, we assume the same from the specific transfer functions of the P&T

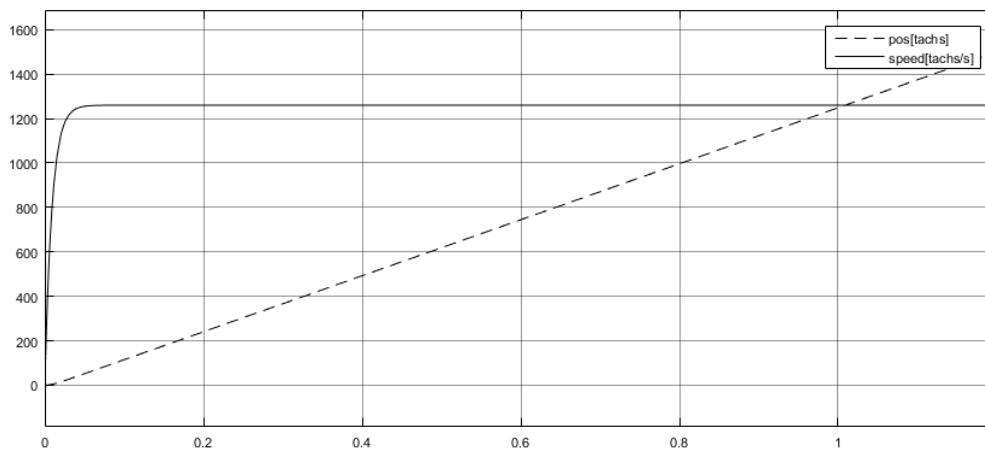


Figure 13: speed & position simulation given a 12V input.

motors.

6.2.2 Tests of the P&T motors

Step response tests were applied to the motors on the actual P&T system. Because of restrictions on the pan part of the system, voltages tested ranged from 3V to 8V. The data has been analysed to find values of constants related to the method of approximating a plant of a DC motor (Figure 14).

Tilt Load/Volt	Measured RPM	DC gain(tach/s/v)	Time constant Tau	Tachs/s
3	34	68,0	0,0649	204
4	51	76,5	0,0713	306
5	69	82,8	0,0671	414
6	85	85,0	0,0688	510
7	105	90,0	0,0707	630
8	125	93,8	0,0761	750
Avg		82,7	0,0698	
Pan Load/Volt	Measured RPM	DC gain(tach/s/v)	Time constant Tau	Tachs/s
3	12	24,0	0,1353	72
4	31	46,5	0,2136	186
5	41	49,2	0,2172	246
6	59	59,0	0,2390	354
7	75	64,3	0,2323	450
8	92	69,0	0,2326	552
Avg		52,0	0,2117	

Figure 14: Table of P&T step response data

As seen below in figure 15, the step response spikes at first. It is assumed this spike occurs primarily because of the following:

In the gear network connecting the actual motor and the pan or tilt part, there are small gaps between individual teeth. The motor is then switched on and while the collective gap closes, the motor runs as if it had no inertia. When the gap is closed the motor is slowed drastically because of the inertia of the pan or tilt part.

The two things together create the spike we see on figure 15 below. Because we assume linearity of the actual systems, we disregard the spike and look at the graphs after the spike has ended to find Tau.

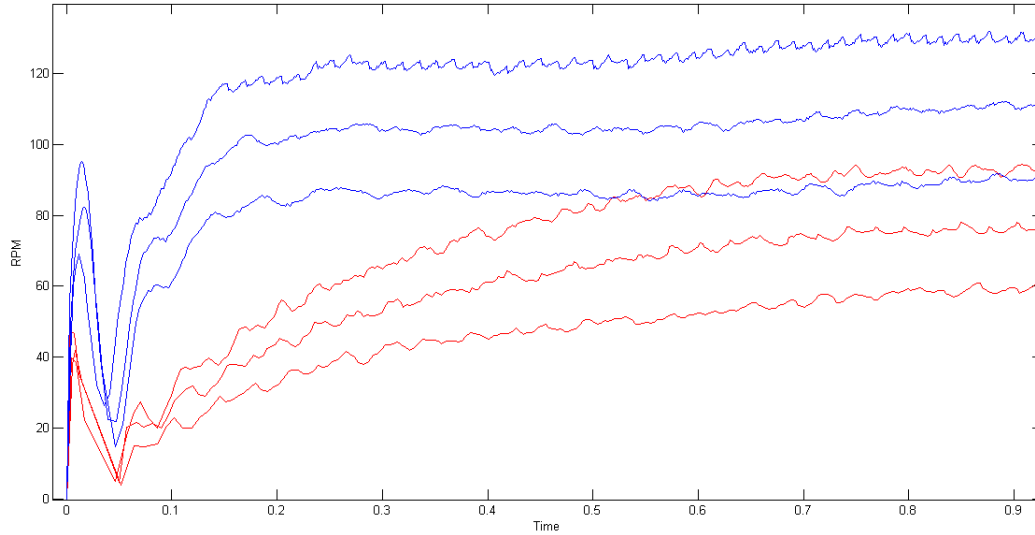


Figure 15: The P&T response to 6V-7V-8V. Blue: Tilt, Red: Pan

Applying the method for creating a plant for a DC motor, discussed earlier, the actual constants for the pan and tilt systems are used in this context to create dedicated transfer functions for each motor in the P&T system. The plants can be seen as equation 8 for the pan part and equation 9 for the tilt part:

$$G_{pan}(s) = \frac{52}{0.212s + 1} \frac{[tachs/sec]}{[V]} \quad (8)$$

$$G_{tilt}(s) = \frac{82.7}{0.07s + 1} \frac{[tachs/sec]}{[V]} \quad (9)$$

6.3 Overview

The mathematical models necessary for predicting and controlling P&T system behavior were found. It was found that a first order approximation was sufficient for modelling the system behavior, because of the negligible armature. Finding proper transfer functions for the P&T systems posed a greater challenge than for the DC motor itself, because of unlinear behavior associated with the gear network in the system. Should an attempt be made to improve the models in the future, this would be an obvious place to begin. These models could be improved either through new tests that try to minimize the unlinearity, or through improvements in the analysis of the data.

7 Control Theory

7.1 Introduction

A controller is an essential part of most autonomous systems. This project is no exception - having control over where the pan and tilt system is pointing at is crucial to meeting the specified requirements.

When discussing what controller to use, figure 16 is the point of reference, where the plant is the previously estimated transfer function, of the DC motor used in the system:

$$G(s) = \frac{105}{0.009s^2 + s} \frac{[tachs/sec]}{[V]} \quad (10)$$

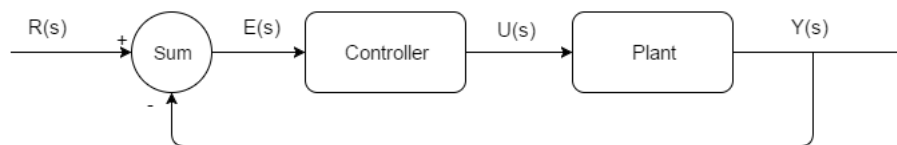


Figure 16: A generic way of representing a closed-loop control system

7.2 Controller selection

For the sake of simplicity Proportional (P), Proportional-Integral (PI), Proportional-Derivative (PD) and Proportional-Integral-Derivative (PID) was the only controllers taken into consideration for this project. Lead-Lag compensators could be considered as well, since in essence they can have similar functionality.

A PID controller is made up by three parts: the proportional gain looks at the current error, the integral looks at the past errors and the derivative looks at current rate of change. Each term has a tuneable gain called the K constants. In general, these constants are what defines a PID controller's behavior. See figure 17 for illustration.




Figure 17: A representation of the PID controller principle

This is the equation of the standard PID controller in the time domain:

$$K_p e(t) + K_i \int_0^t e(t) dt + K_d \frac{de(t)}{dt} \quad (11)$$

If transformed to the s domain the equations looks like this:

$$G(s) = K_p + \frac{K_i}{s} + K_d s = \frac{K_p s + K_i + K_d s^2}{s} \quad (12)$$

7.2.1 Proportional Controller

The P controller is by far the most straightforward controller, being that it is simply adding a gain to the open-loop transfer function. A P controller is a desired controller, since it is very simple to implement and tune, though this project have to meet the stated requirements.

To help analyse and give some intuition about how the motor's transfer function behaves, a root locus of the open-loop transfer function is plotted see figure 22. As one would expect it seems like this kind of controller will not suffice for the project.

As seen in figure 18, the step response of the closed-loop transfer function, has an overshoot of 14.1%, and a settling time of 0.055 seconds. Meeting the specified requirements for the project with this controller is not possible. The overshoot can be decreased by reducing the gain below 1, but this induces a longer settling time and vice-versa.

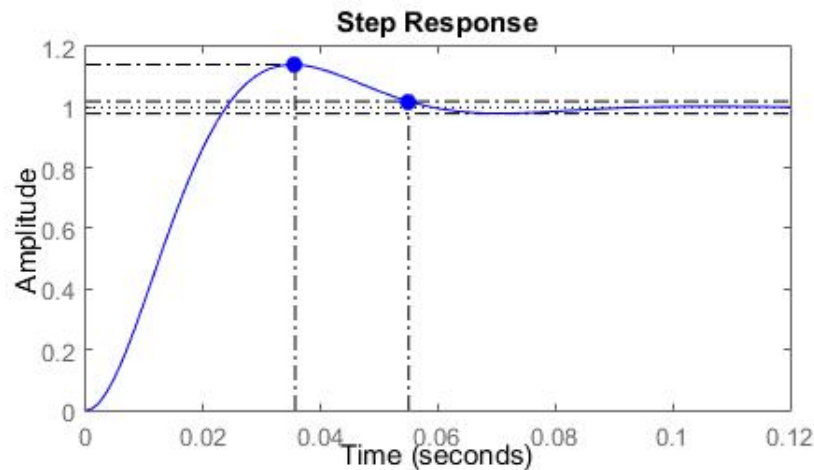


Figure 18: Peak amplitude: 1.14 - Overshoot(%): 14.1 - At time(seconds): 0.0356 Settling time (seconds): 0.055

7.2.2 Proportional-Integral Controller

The PI controller accumulates past error terms over time to eliminate steady-state errors in the system. This can cause an increase in overshoot called the integrator wind-up, and an increase in settling time. Though if a requirement is complying with a steady-state error of 0, then it is up to the designer to decide, if reaching steady-state is more vital to the system, than having more overshoot and a longer settling time.

Using a PI controller is the same as adding a pole and a zero to the open-loop transfer function. As seen in figure 22, this controller does not seem to live up to the requirements either. Placing the zero and pole differently could make a better controller, but it would never be able to adjust to both design criteria.

The step response of the PI controller is seen on figure 19, the only difference from the P controller is a slightly longer peak time and settling time and a larger overshoot, as expected. This controller would not meet the requirements.

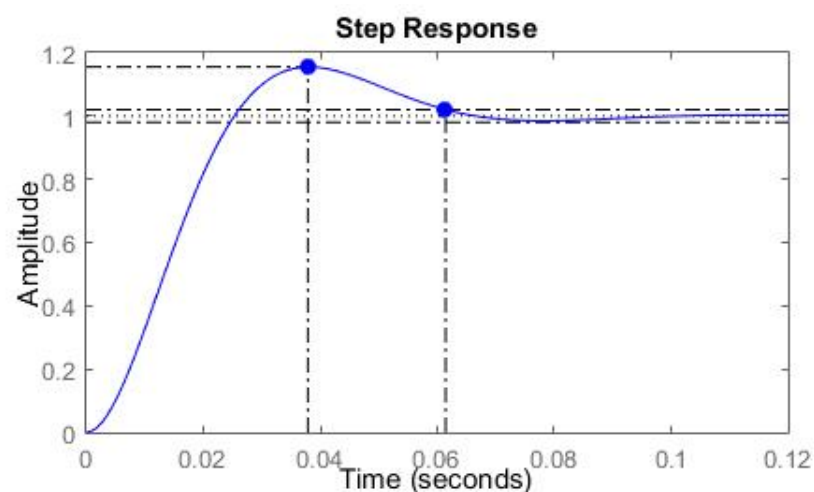


Figure 19: Peak amplitude: 1.15 - Overshoot(%): 15.5 - At time(seconds): 0.0379 Settling time (seconds): 0.0614

7.2.3 Proportional-Derivative Controller

A PD controller is equivalent to the addition of a simple zero, which improves the transient response. From a different point of view, the PD controller may also be used to improve the settling time or stability, because it anticipates large errors and attempts corrective action before they occur.

Figure 22 shows a promising root locus of the PD controller, it seems like that with the right amount of gain the system would fulfill the requirements.

The step response on figure 20 shows, that in fact this controller would be ideal for this project. There is no overshoot, and the settling time is well below 0.05 seconds.

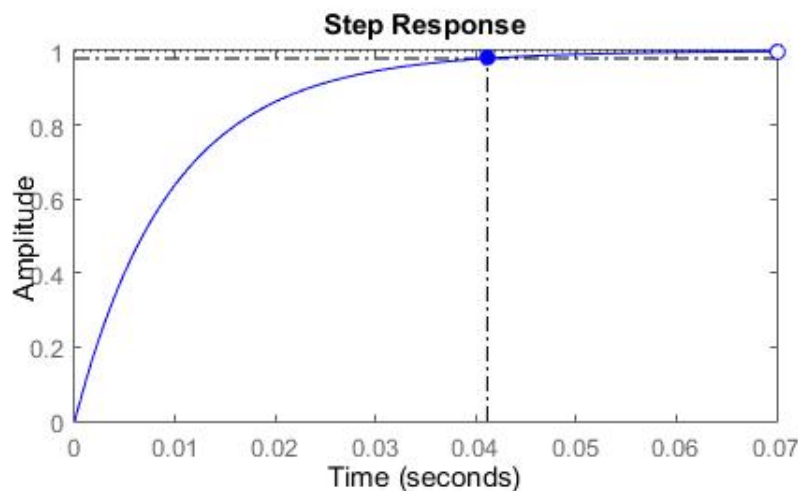


Figure 20: Peak amplitude: ≥ 0.998 - Overshoot(%): 0 - At time(seconds): > 0.07 Settling time (seconds): 0.0411

7.2.4 Proportional-Integral-Derivative Controller

The PID controller is the most complex controller which is going to be discussed. As noted earlier a PD controller was sufficient for generating the wanted response. So why discuss the PID? The reason is that real world systems might have limitations that the analysis have not taken into account yet. Reaching a steady-state error of 0 is actually a problem with a PD controller for this project, since the pan and tilt system does not function below a certain range of voltage, which means that an integral is needed.

Notice on figure 22, that a PID controller is the addition of 2 zeros and 1 pole to the open-loop transfer function. The root locus shows that this system with a relative gain still meets the requirements of the design.

The step response of the PID controller, figure 21, shows that indeed a PID controller would be a good choice of controller for this project. There is well below 10% overshoot and the settling time is a great deal below 0.05 seconds.

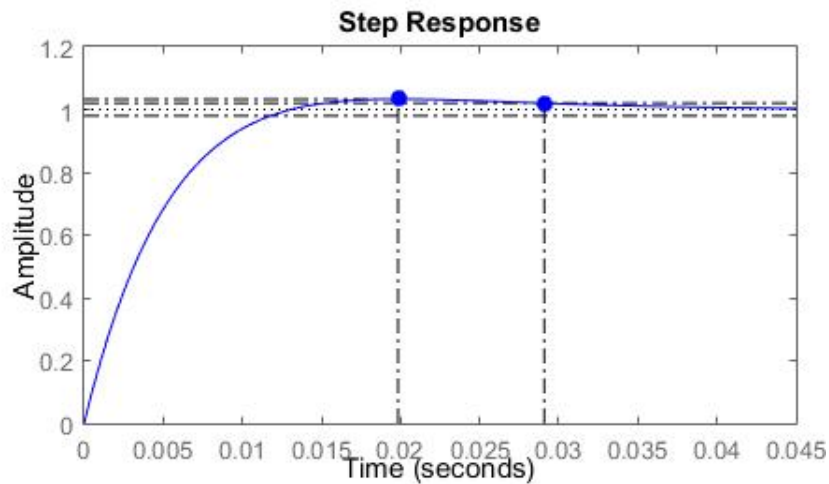


Figure 21: Peak amplitude: 1.03 - Overshoot(%): 2.92 - At time(seconds): 0.0201 Settling time (seconds): 0.0277

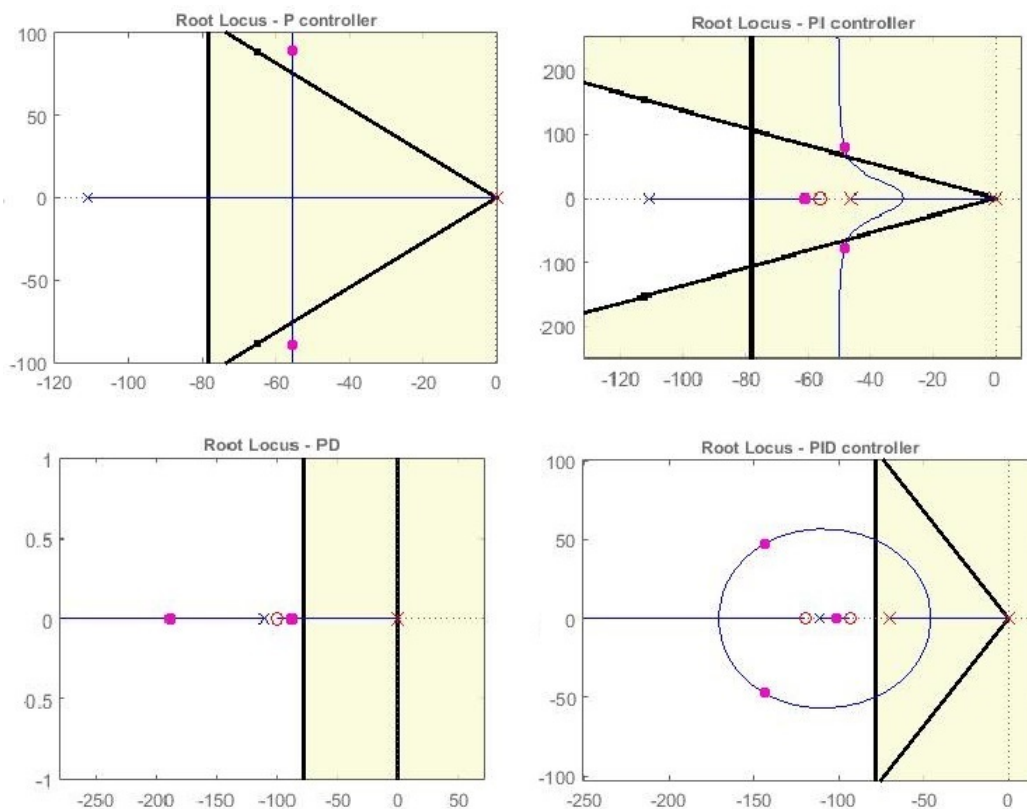


Figure 22: Open-loop transfer function Root Locus plots of the estimated motor transfer function, with different controllers. The shaded areas indicates unsatisfied design specifications, if poles of the closed-loop transfer function are placed there. The designs are: Overshoot > 10%, settling time < 0.05 seconds.

7.3 K constants

Once either a P, PI, PD or PID controller has been chosen for the project, the next step would be the estimation of K-constants for the respective controller. Finding these constants to meet one's requirements can be done in various ways. In this section the Manual tuning and Root Locus method will be discussed.

7.3.1 Manual tuning

This method requires the system to be online and running for tuning, though if a simulation of the system has been made this tuning method is still viable on the simulation - the catch is that a simulation's artificiality cannot be denied, thus the actual online system proves more truthful.

There is no right way to tune manually, since experience can lead to a quick tuning. A known procedure is, however, to first set one's k_i and k_d to zero, and then increase the k_p until the system starts to oscillate, k_p should then be set to approximately half this value. The k_i is now increased to correct for any offset within the desired time. Lastly, if needed, the k_d is increased until the requirements are met.

7.3.2 Root Locus tuning method

This method is a mathematical approach with the help of software like MatLab to find k_p , k_i and k_d from the root locus. It is not required to have the online system up and running to use this method, which in many cases can prove useful (imagine a system where unpredicted behavior can have fatal consequences). Though a mathematical model of the system is needed.

The root loci on figure 22 can be used, since they are the open-loop transfer function of the system. It has already been established that using a PID controller would be a good choice to meet the project's requirements. Therefore this example is based on the PID root locus from figure 22. The root locus was made with the SISO tool from Matlab with the open-loop transfer function equation 7. The two zeros at -90 and -55 on the real axis and a zero at the origin is what makes up the PID controller, and they were placed there to satisfy the design requirements. The SISO tool Compensator Editor of this root locus shows these zeros and pole as a compensator, which is another way to represent a PID controller. In the example of the PID controller from figure 22 the compensator looks like this:

$$200 \frac{(1 + 0.011s)(1 + 0.019s)}{s} \quad (13)$$

Rewriting equation 13 to look like equation 12 makes it quite easy to read the K constants.

$$\frac{0.0396s^2 + 5.8s + 200}{s} \quad (14)$$

As seen in 14 the K constants for the PID controller are $K_P = 5.8$, $K_I = 200$ and $K_D = 0.0396$.

7.3.3 Controller constants for Pan and Tilt systems

Applying the root locus method to the transfer functions for the Pan and Tilt systems, PD and PID controller constants were derived for both of these. These constants can be found in figure 23.

Tilt	PD Controller	PID Controller	Pan	PD Controller	PID Controller
Zero 1	-10	-5		-2	-3
Zero 2	-	-40		-	-20
Pole	-	0		-	0
Kp	1,3	13,5		1,1	38
Ki	-	60		-	100
Kd	0,13	00.03		0,55	1,65
Settling Time	0,0474 s	0,0447 s		0,0472s	0,0466s
Overshoot	~0%	5.74%		~0%	3,54%

Figure 23: Table with K constants for PD & PID controllers for the Pan- and Tilt systems

7.4 Simulation

Making a simulation of the system is a good way to figure out whether the controller needs to be redesigned or tweaks needs to taken into consideration. Since the motors on the pan and tilt are expected to have similar characteristic as the load free motor (though pan and tilt are expected to have different time constants), the simulation is then first done on the load free motor transfer function with the PD controller from figure 22. The K constants were found with the same method as discussed in the Root Locus tuning Method section. The constants are $K_P = 2$ and $K_D = 0.02$. The simulation is done in MatLab's Simulink software, see figure 24. The Step block is a step input of 1, and since the DC gain of the plant is in the units of $\frac{tach}{sV}$, this simulation illustrates the system rotating the motor 1 *tach*. The step input is fed to the feedback summation, where the feedback is subtracted from the reference input. The summation result is now the error signal, and this signal is then fed into the PD controller. The PD controller reacts depending on the rate of change of the error signal and passes on its calculated value. The motor plant produces physical rotation, which is integrated to *tach.s*. The output is now sampled as feedback and is then fed back to the reference signal. The Time Scope will generate the plot of the simulation.

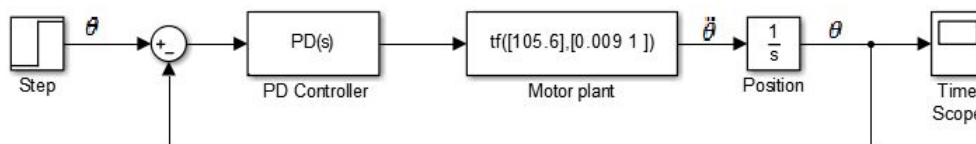


Figure 24: A simulink simulation of a PD controller of the load free motor's transfer function.

The step response of the simulation is expected to look like the step response from the PD analysis. As seen on figure 25, it looks a lot like expected. Upon inspecting the time 0.0226 seconds, the amplitude is 0.982 which is within 2% of the settling time, so the simulation is very close to the expected.

Since this is a simulation of the physical system, the physical aspects that are obvious should be taken into consideration. The motor of the system cannot take inputs higher than 12 volts, therefore a saturation limit will be added to the PD controller to simulate this. Other than that, the friction and inertia in the system makes it impossible for the system to move below 2 volts. This means that inputs ranging from -2 volts to +2 volts to the plant, must always be set to 0 to simulate the idling of the motor. On figure 26 these changes can be seen.

The step response of figure 26 is seen on figure 27, and as it is seen that the motor stops before reaching the target. This means in reality, the motor never reaches the target, and since this is a secondary requirement this is

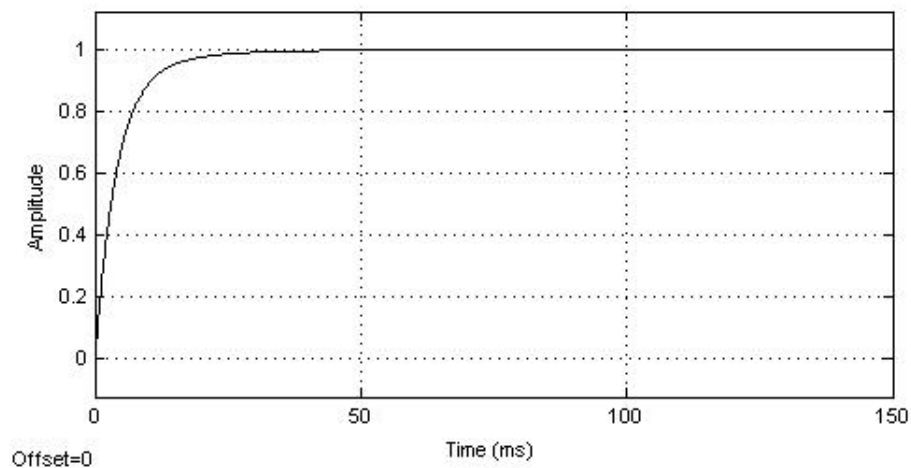


Figure 25: The step response of the simulation from figure 24




Figure 26: A Simulink simulation of a PD controller of the load free motor's transfer function with saturation of -12 and +12 voltage and constraints with output 0 at voltage range -2 to +2 volts .

something that would be nice to compensate for. Luckily a PID controller has been discussed, and the constants for it is already known, and this controller is ready for simulation. Figure 28 is the step response of the simulation of the PID controller, with the constants found in the Root Locus tuning method section, in the same setup as figure 26. The step response shows that with a PID controller the simulation now reaches the desired target while still having a settling time below 50 ms. In conclusion a controller for the load free motor has been designed and simulated to work for the desired requirements.

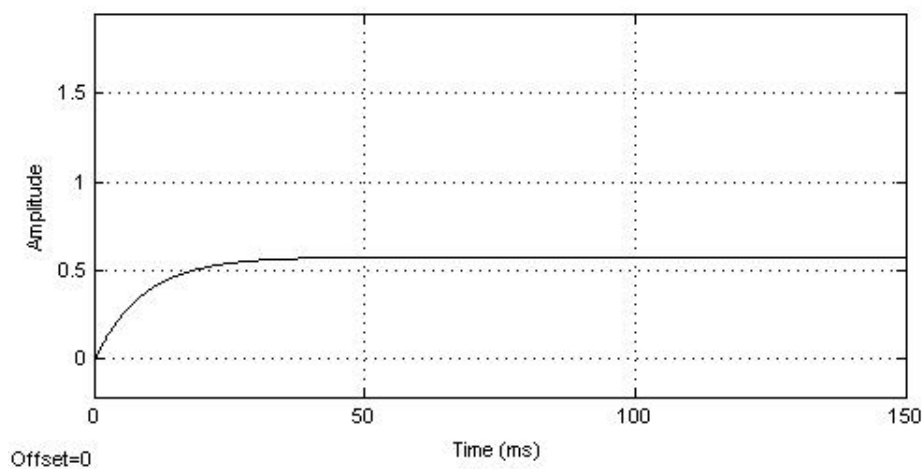


Figure 27: The step response of the simulation from figure 26 with a PD controller

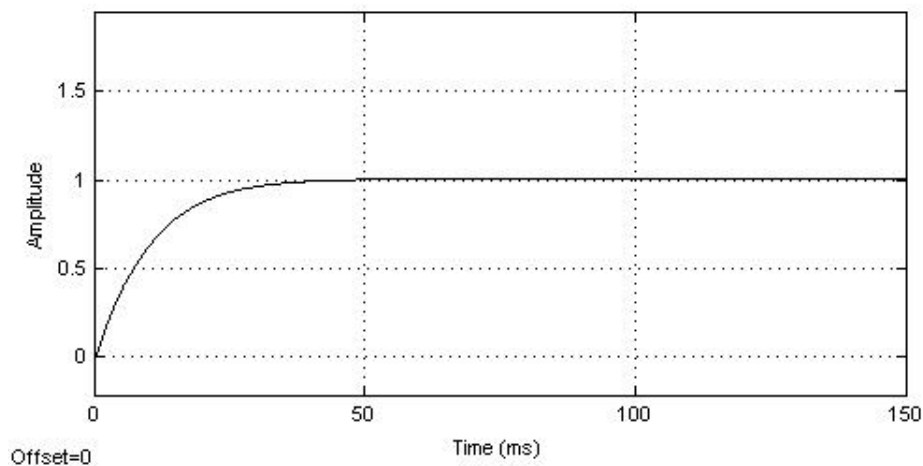


Figure 28: The step of the simulation from 26 with a PID controller

Since the simulation of the load free motor needed a PID controller to meet the requirements, the pan and tilt motor on the actual system is expected to be of the same behavior, and would also need a PID controller. Also the system constraints on the system might prove of a problem for the PID controllers of the pan and tilt, since the saturation and the deadband between -2 volts and +2 volts, has not been take into considerations while designing the control system. The transfer functions of the pan and tilt were simulated in the same manner as in figure 26 with both PD and PID controller constants shown in figure 23.

On figure 30 the step responses of the pan motor with PID and PD control can be seen, it is clear that, as expected, the PD control does not reach an acceptable steady state, while also being a lot slower than the required settling time of 50 milliseconds. The PID control however reaches an acceptable steady state. Nonetheless the settling time of 0.689 seconds is also too slow, and there is an overshoot of 11.2%.

The slow settling time and overshoot can be explained by the saturation of the plant, and the deadband between -2 volts and +2 volts. The designed controllers do not compensate for these factors.

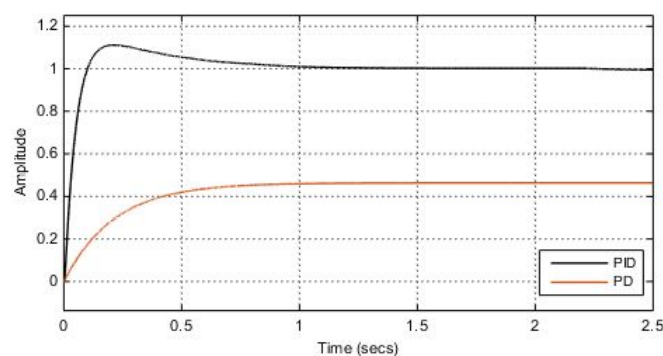


Figure 29: Step response of the simulation from figure 26 with pan motor and results of both PD and PID controllers.

Likewise on figure 30, the step responses of the tilt motor with PID and PD control can be seen. As expected the PD control never reaches an acceptable steady state, and the settling time is slower. The PID controller reaches an acceptable steady state, but also have a long settling time of 0.491, while also having an overshoot of 18.7%. Same

as for tilt motor, these behaviors can be explained by the saturation of the plant and the deadband between -2 volts and +2 volts.




Figure 30: Step response of the simulation from figure 26 with tilt motor and results of both PD and PID controllers.

Another simulation on the pan and tilt motors were run. This time an step input of 540 was set. This is equal to half a revolution on either motors. This simulation was run, because in practice, the system should be able to handle the kind of inputs. On figure 31 the step response of 540 can be seen, as noted the pan and motors overshoot with about 100% and the settling time is greatly increased because of this. This happens because of integrator windup, and was not compensated for when designing the controllers. In the simulation an anti windup can be set, to simulate an implementation of windup protection. On figure 32 the result of a step input of 540 with windup protection on can be seen. The pan and tilt on figure 32 seems to behave acceptable, the tilt has no overshoot and the pan has an overshoot of 3.4 % overshoot. Though the settling time is still slow, however this is to be expected as the system is saturated.




Figure 31: Step response of the simulation from figure 26 and 540 input instead of 1. The pan and tilt motor are without wind up protection.

The transfer functions and the controllers for the pan and tilt have now been tested in simulation, and the designed controllers seems promising. The tilt motor would on a small step input not fulfill the requirements of less than 10% overshoot and the settling time requirement is not met either. It would seem like that reaching a settling time of 0.05 seconds or less would be impossible for the tilt, since the system cannot move fast enough. Tough for a step input of a larger scale, the overshoot is high zero and well within the requirement. As for the pan motor, a small step input has an overshoot really close to 10%, but a long settling time. The requirement of having overshoot less than 10% is therefore not fulfilled, either is the having a settling time of 0.05 seconds for the pan motor. On a larger input the pan motor has an overshoot of only 3.4% and is well within the requirement.

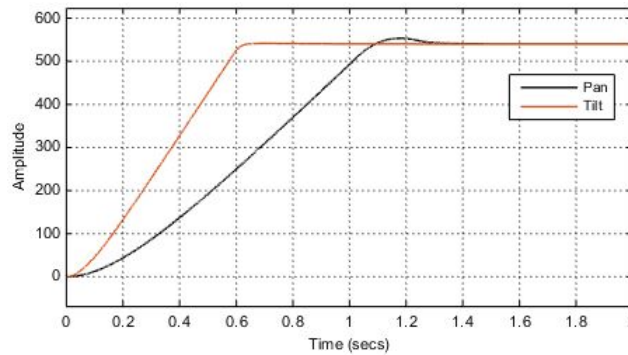


Figure 32: Step response of the simulation from figure 26 and 540 input instead of 1. The pan and tilt motor are with wind up protection on.

In general, the overshoot requirement of 10% seems to met for step inputs of a relative size, and the settling time requirement of 50 ms is impossible for the system to meet, because of saturation and deadband.

7.5 Implementation - digital

7.5.1 Direct implementation

One way of implementing the controller is a direct implementation of the PID-controller. In this case, the plant input will be a summation of the 3 PID contributions, that will all have to be calculated in real-time. This method makes a lot of sense intuitively, and provides easy access to modifying the K-constants during the tuning-process, as these are simply constants that are multiplied with the 3 contributions. In order to use this method, the mathematical operations of integration and differentiation will be have to implemented.

7.5.2 Filter approach

An alternative way to implement the PID controller is converting it into a discrete-time filter. This method has the advantage of not having to do any actual integration or differentiation on the FPGA. Given a continuous-time transfer function and a sampling rate, the matlab command “c2d” can convert the transfer function to the z-domain. Through algebraically isolating the current output of the controller, the filter coefficients for the implementation can be calculated.

The sampling rate should ideally be as high as possible, as this will improve the accuracy of the discrete filter. Through matlab tests, it was found that sampling rates of higher than 10kHz do not create further improvement in the accuracy of the filter, and thus there was no need to increase the sampling rate beyond this value, even though it was technically possible.

7.5.3 Discrete filter for Pan-system [10kHz]

Using the above method, discrete filters were calculated for the Pan- and tilt systems.

$$\frac{u(z)}{e(z)} = \frac{1.65z^2 - 3.3z + 1.65}{z^2 - 2z + 1} = \frac{1.65 - 3.3z^{-1} + 1.65z^{-2}}{1 - 2z^{-1} + z^{-2}} \quad (15)$$

$$u(z)(1 - 2z^{-1} + z^{-2}) = e(z)(1.65 - 3.3z^{-1} + 1.65z^{-2}) \quad (16)$$

$$u[n] - 2u[n-1] + u[n-2] = e[n] - 3.3e[n-1] + 1.65e[n-2] \quad (17)$$

$$u[n] = e[n] - 3.3e[n-1] + 1.65e[n-2] + 2u[n-1] - u[n-2] \quad (18)$$

7.5.4 Discrete filter for Tilt-system [10kHz]

$$\frac{u(z)}{e(z)} = \frac{0.3z^2 - 0.5999z + 0.2999}{z^2 - 2z + 1} = \frac{0.3 - 0.5999z^{-1} + 0.299z^{-2}}{1 - 2z^{-1} + z^{-2}} \Rightarrow \quad (19)$$

$$u(z)(1 - 2z^{-1} + z^{-2}) = e(z)(0.3 - 0.5999z^{-1} + 0.299z^{-2}) \quad (20)$$

$$u[n] - 2u[n-1] + u[n-2] = 0.3e[n] - 0.5999e[n-1] + 0.299e[n-2] \quad (21)$$

$$u[n] = 0.3e[n] - 0.5999e[n-1] + 0.299e[n-2] + 2u[n-1] - u[n-2] \quad (22)$$

7.6 Overview

The necessary coefficients and constants have been found to complete the implementation-method of choice. The discrete filter method, however, was discovered quite late in the chronology of the project, and due to this lacks verification and testing. Most of the project work has been carried out under the assumption that the direct implementation method was the best, and thus this method has been given the main focus during testing and implementation.

8 The Microcontroller

Given the design requirements in section 3.2, and the block diagram in section 4, all the functionality that is to be placed in the microcontroller is defined. From here the next stage is to define the content of these blocks and how they interact with each other.

The microcontroller is to establish an application programmer interface (API) to the P&T driver placed in the FPGA. The microcontroller will also run the application that uses this interface in order to help the user achieve his/her goal, which is to control the spotlight.

The program will be written in the C language and it will be documented using task diagrams and state machines.

8.1 Choice of Operating System

It is safe to assume that because this program will be using more than one task, some sort of operating system (OS) is required. The decision of which operating system to use, will have great effect on the entire project, since it defines the domain in which the entire application will have to operate in. This OS must therefore as a minimum support the ARM Cortex M-4 processor, which is used in this project.

By using a Real Time Operating System (RTOS) like freeRTOS, it would give the advantages of preemptive scheduling. Being able to preempt a low priority task to pave the way for a task of higher priority, is an advantage when low response time is important. However, the application will only benefit from using an RTOS if its size extends 1 MB. Application smaller than 64KB simply would not be large enough for an RTOS to have any effect relative to a standard OS.⁵

Assuming that the project will reach this size, combined with the fact that this OS is used in one of the courses during this semester, freeRTOS was chosen.

8.2 Program Structure

After the definition of the operating system, the next stage was to define the actual application that needs to be implemented. Given the block diagram of figure 2, the blocks that make up the system were already defined.

8.2.1 Defining tasks

First step is to identify the tasks that make up the system. A good rule of thumb is to make a driver task for every piece of hardware that the system is to interact with. This idea is based on the Parallel criteria, which implies that things can be separated into different tasks if the functionality run independently and simultaneously, which indeed will be the case for a hardware driver. However the UART0 driver that supports connection between the PC and micro controller, can send and receive data independently. Therefore this task is split up in to two tasks. A receive task, and a transmit task. Doing so also have the advantage of reusability, for other applications requiring the same hardware.

Now when all the hardware is taken care of, the only thing remaining is the application itself. The application should be able to handle commands given by the user. Commands could be "set coordinate pan, tilt", "run light-show #2" and "set minimum velocity tilt". Since the entire application will consist of function calls, the main body of the application will take the shape of a kernel task that executes the functions defined in the function list shown in appendix A. However by executing the command "run lightshow" it requires of the application to feed the P&T driver with new target points once the previous target point has been reached. Considering that one still would like to send commands while running a light-show this will make two tasks based on the parallel criteria. The kernel Task and the light-show task.

Last but not least the application operates on the current position in order to handle a light-show. But gaining the current position requires pulling from the application. This task therefore makes the final task in the system based on the Periodical Criteria.

⁵Texas Instruments, Why use RTOS in microcontroller Applications, <http://www.ti.com.cn/cn/lit/wp/spry238/spry238.pdf>, 23/05/16

The application must therefore consist of the following tasks:

Driver tasks:

- UART0_rx_task
- UART0_tx_task
- Joystick_task
- SPI_Master_task

Application tasks:

- app_kernel_task
- app_lightshow_task
- app_update_current_task

8.2.2 Connecting the tasks

Now the tasks are to be linked together. This will happen in a structure so that every task that requires input from other tasks has one and only one queue to pull from. This is decided since the OS can pause the task while waiting for one queue. The UART tasks and the joystick tasks will both send commands to kernel via the application queue, and a shared state memory. The same goes for the connection between the kernel task and the light-show task.

Each queue needs to be semaphore protected in order to prevent a race condition when reading and writing to the queue, but since freeRTOS handles who gets to use the queues, it must be safe to assume that such conditions are taken care of in the OS.

The shared state memory (SSM), needs semaphore protection as well. Since the number of memory locations is relatively small (less than 64), combined with the fact that the critical section only consists of one read or write, the entire SSM will be handled by the same semaphore for simplicity.

The two application tasks `app_update_current_task` and `app_kernel_task` will both be accessing the `spi_tx_queue`, but this will only be a problem if they wish to send multiple messages in context, which is not the case for this application. The same goes for the `application_queue`. Here the send messages and events content will be independent from each other.

Applying all these thoughts will result in the final application task diagram as shown in figure 33. Now when the task diagram has been made, it is time to consider possible structural weaknesses such as deadlocks and handling of queues.

8.3 The Application Programming Interface

When handling a piece of hardware like the P&T system or the UART0 unit, it is favorable to make reusable drivers, but if the unit is to interact with the application following specific protocols, you quickly lose the reusability. Instead you may be writing a separate module called an Application Programming Interface (API), containing all the features necessary, for the application to handle the driver easily and correctly.

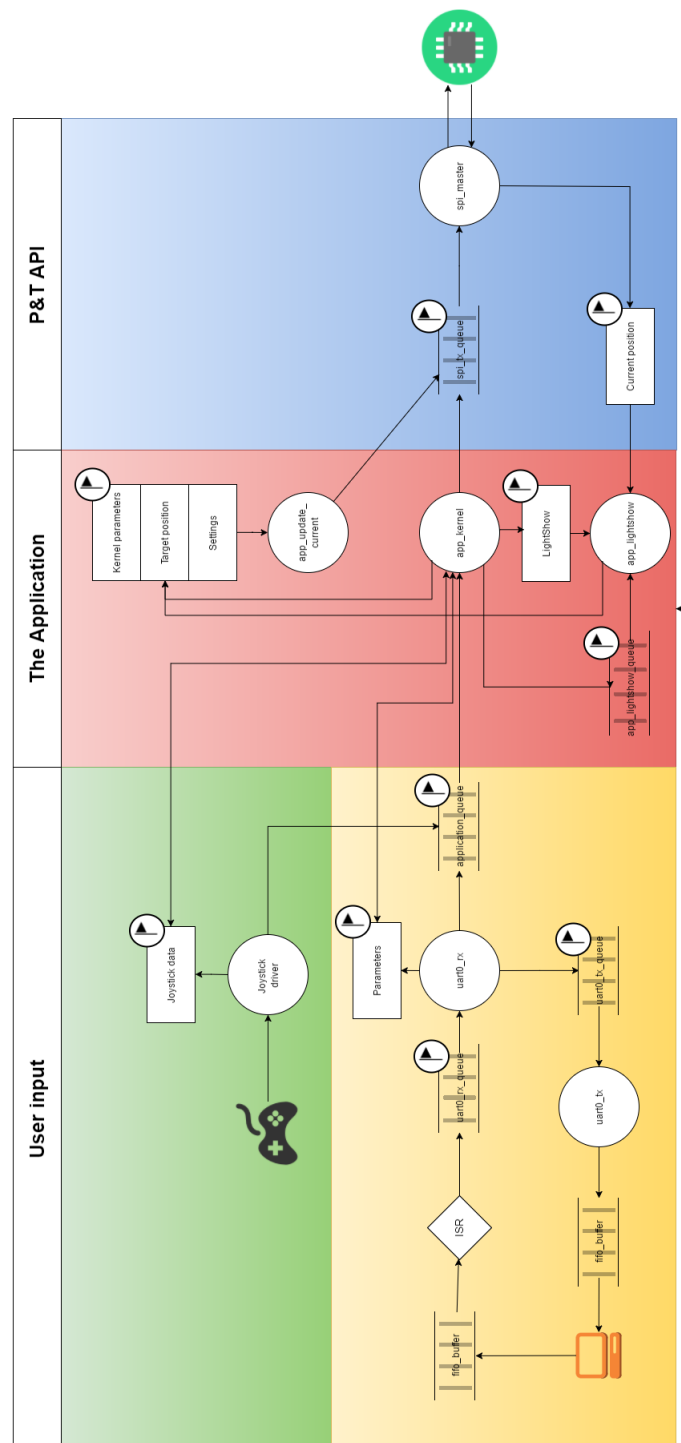


Figure 33: The application task diagram

An example of this can be found in the P&T API. This module handles the format of the messages sent to the `spi_tx_queue` and received by the `spi_master_task`. Here the message must have "n" address bits and "m" message bits, but say that we one day were to change this message format to something different, then it would only effect the P&T API module, while leaving the rest of the application as is. Using API's will therefore make the application more agile.

8.4 Event Queues and Shared State Memory

When operating with tasks, it should come as no surprise that it will be implemented as a state machine. This structure will operate in one state until a specific condition is met, from here the tasks will operate in ways defined by the new state. For this application the tasks will mostly switch states as they receive the correct events. In short an event is a number that the programmer has assigned a specific meaning to. It is therefore important for the application to know if it is reading an event or not. In other words, if the sending task uses the queue for both user data and events, it will be impossible for the receiver to know, if the number 214 is to be interpreted as the value, or the event e.g "UPDATE_DISPLAY".

Suppose an event is an eight bit unsigned integer, and it is sent to a task through a queue. It is impossible for the receiver task to know what data type it is, and who sent it. It will therefore have to send extra information stating the data type in the shape of a header message. An alternative would be splitting up the communication in a `event_queue` and a SSM buffer. Only events will go in the queue, and data will first be placed in the SSM and then an `SSM_READ` event is send to the receiver task. From here the receiver will read the data and set it to zero. This structure is shown in figure 34.

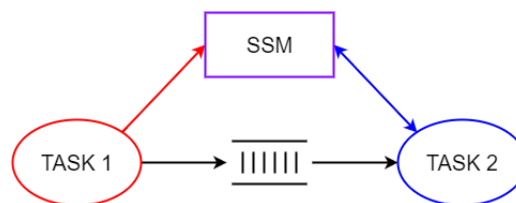


Figure 34: Communication using event queues

It is possible to first send a header message telling what to do with the next message, however, this solution would not be particularly robust, since all it takes for the system to get lost is for one message to fail at either getting sent or received. It will then go out of synchronization and the header will be read as data and the other way around, and since this error only happens rarely, this error will be hard to debug. It is therefore preferred to keep the messages atomic, by using a queue and a SSM address.

8.5 Deadlock Protection

A deadlock has happened when two or more processes wait for each other to release a resource, but since they all are waiting, no one will ever finish their task. The tasks and resources will therefore be unusable or deadlocked for the rest of the applications lifetime.

1. **Mutual Exclusion:** is the concept of making sure that two or more concurrent processes are not in their critical section at the same time.

2. **Hold and Wait:** is when a process holds on to a resources and waits for the rest of the resources it needs.
3. **No Preemption:** means that only the process that holds the resource can release it. If preemption is allowed, the program is able to take resources from a process and allocate it to another process if needed.
4. **Circular Wait:** is when one process is waiting for at resource that another process has, and the new process is waiting for a resource held by the first one. As long at this makes up a closed chain of waiting processes, there is a circular wait.

All four conditions need to hold true in order for a deadlock to occur. You can therefore provide a deadlock-free environment, by avoiding one of these conditions⁶.

To avoid deadlocks in the system, the circular wait condition is removed via semaphores. However since the system only has one semaphore, circular wait will never be an option.

8.6 SPI-Master

The microcontroller has to share information with the FPGA, and according to the requirements mentioned in section 3.2.1, this communication has to happen using Serial Peripheral Interface (SPI). The protocol will operate as described in section 5. It is therefore important that this link never will be a bottleneck for the system application.

The minimum bit-rate requirement depends on what data that needs to be send, and how often it needs to send it. Since the PID-controller is placed in the FPGA, the application only needs to update the current position once every 5 ms, which is the minimal delay that freeRTOS can put on a task.

The length of the messages sent is defined by the driver placed in the FPGA. As described in section 10.2 all messages will have the length of 14 bits. Updating both pan and tilt requires two messages. Considering all other messages but the main pulling sequence as negligible results in a total bitrate of:

$$Bitrate = \frac{14bit * 2}{0.05s} = 5.600bit/sec \quad (23)$$

The Cortex M4 has a built-in freescale spi module, that supports four different variations of SPI, see appendix ?? . However, implementing the SPI from scratch, gives the highest degree of control over the process, which is preferable in this project.

The state machine for the task is illustrated in figure 35. It consists of three states and is initiated waiting for a message to transmit from the spi_tx_queue. When a message is pulled from the queue, it jumps to the send state, and after sending the message, the received message will be handled by the P&T API in the Receive task. From here it will only switch back to the idle state if the API call was successful, else it will continue until it succeeds.

8.6.1 Performance Test

After implementing the the SPI-master task in the final application, the max bit-rate was tested by overloading the spi_tx_queue, and giving the SPI-master task High priority in the OS. This made the the task perform at maximum

⁶Operating system Concepts 9th edition, by Abraham Silbershatz, page 319

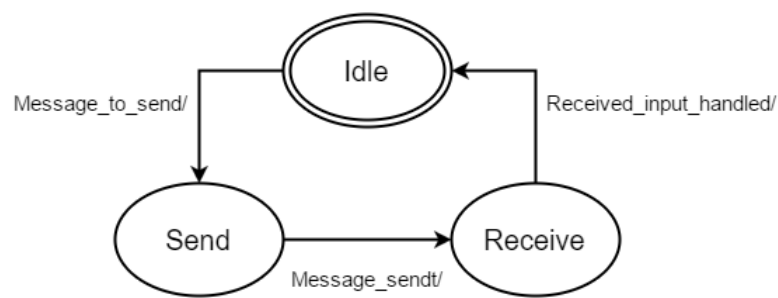


Figure 35: The SPI-master state machine

capacity. The measurements of the SPI connection can be seen in figure 36. Since it is able to transmit messages of 14 bits at 7.7 kHz the bitrate will be the product of these two values giving the result 107 kbit/sec

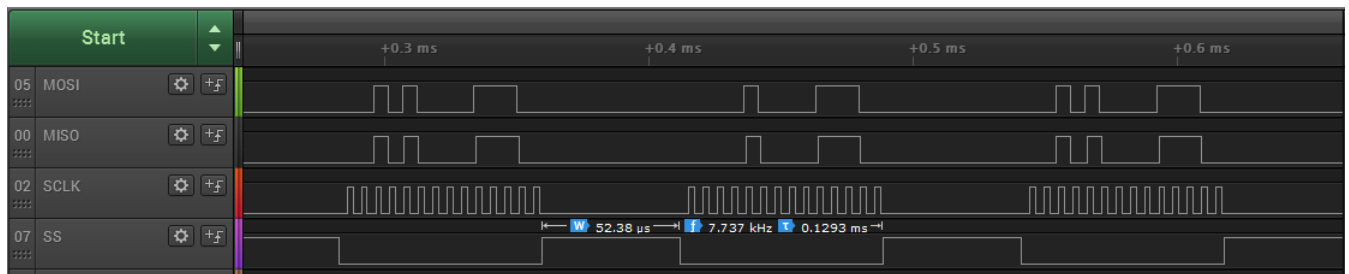


Figure 36: Shows the max performance of the SPI-Master task

8.7 The Application

When the interface for the drivers is set up, the only thing left is to implement the application. The first part in the application to implement is the `app_update_current_task`. The responsibility of this task is to update the SSM holding the current position of both pan and tilt. This will be implemented by the use of one api call into the P&T API.

The next task to implement is the `app_lightshow_task`. This task's responsibility is to feed positions to the P&T driver, doing the light-show. This is implemented by sending a new target position to the driver once it has reached the previous target position.

The final task to implement is the `app_kernel_task`. This task has to handle the commands and parameters given by the user. First, the commands are separated into three groups, depending on the number of parameters it requires. After a command is received the kernel waits for the user to enter the parameters needed for the kernel to execute the command. This structure is implemented in a state machine as shown in figure 37.

8.8 Application Conclusion

The functionality described in this report represents the fully-dressed system containing all the features that would be needed while working with a real spotlight. However, due to time constraints, some of the secondary requirements were cut from the final product.

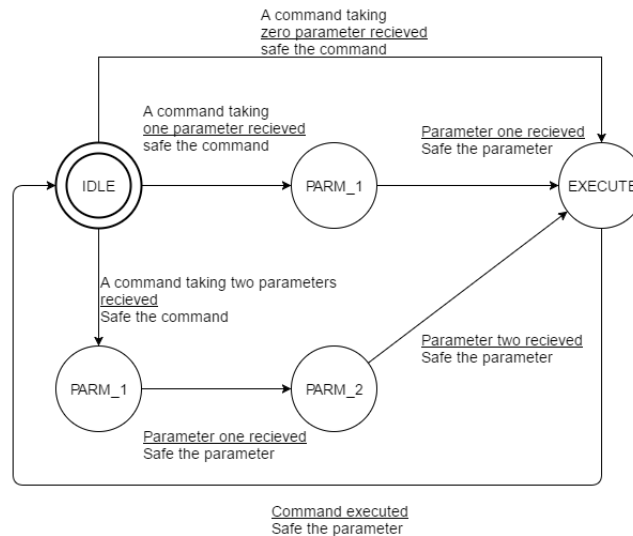


Figure 37: The kernel state machine

The joystick: Being able to control the spotlight via a PS2 controller is a secondary requirement. A working driver was made but the corresponding API was not. This meant that the feature was not implemented. This is not a problem, as it is possible to control the system via the UART connection to the PC.

Position vs given coordinates: The input from the P&T system is tachs(falling and rising edges from the Hall effect sensors). This corresponds to an orientation on a sphere. Since the target for the spotlight is a flat surface, some conversions are needed in order to translate the desired position on the scene to an amount of tachs. The translation takes place in P&T API, and must handle conversion in both directions.

Following a walking person: The kernel is built with the option to change the speed, based on user input, but the only way to change this dynamically and fast enough will be through the analog input of the joystick. But since the joystick was never implemented, tracking a person was not an option.

9 Driver in the FPGA

This section describes how the driver is implemented in the FPGA. The section starts with an overall analysis of the system and afterwards goes into details with some blocks defined in the analysis and ends with a discussion on how the PID controller could have been implemented better as a filter.

9.1 Overall analysis

The easiest way to describe how the overall system works is by looking at the general components and how they are connected. The general components used in the implementation are: a comparator, a controller, a sensor, a motor driver with a PWM module and a SPI slave with a data bank. A generic multiplexer display was also added to simplify the debugging process, however this should not be seen as a part of the final product. Since the goal is to control two motors, a separate controller setup was made for each motor. This means that each motor is connected to its own setup of a comparator, a sensor, a controller and a motor driver. The SPI slave component is common for the motors. Figure 38 is an illustration of how the components are internally connected. This figure

also shows information on the signals between the individual components.



Figure 38: Placeholder figure

The components can be divided into the following blocks, each with a specific purpose.

The first block is the SPI slave block which contains only the SPI slave component. The purpose of this block is to receive and store information used by the rest of the components. It then shares the information with the components that are in need of them. An example of this could be the maximum and minimum speed of a motor which is sent to the controller (this connection can also be seen on figure 38).

The second block is the controller block which contains the comparator and the controller components. The purpose of this block is to regulate the motors according to a target input and the feedback. The block generates a duty cycle and a direction, that is passed onto the motor driver.

The third block is the sensor block which consists only of the sensor component. The purpose of this block is to keep track of the current position of the motor. This is done by looking at the hall sensors from the motors and relating their output to the current position.

The last block is the motor driver which contain the motor driver component and its internal PWM module. The purpose of this block is to control and power the motor. This means that the block generates a PWM signal and outputs it on the correct pin according to the desired direction.

10 SPI slave block

This section will describe in detail how the SPI slave in the FPGA receives and transmits data, according to the description of the SPI communication protocol mentioned before.

10.1 Description

The SPI slave has two different tasks. The first task is to exchange data with the SPI master in the microcontroller. The second task is to store the information that is received and present it to the components that need the data. The SPI slave could be called also some sort of information bank for the driver.

The data exchange part follow the SPI communication protocol mentioned in section 5. The information bank part is made by having the block hold the information in the actual output of the block. This is done so that the data is constantly supplied to the block in need of the data. However this requires a lot of data paths, but this is not much of a problem since this is going on internally in the FPGA. The reason the data is stored in the SPI slave component was also to make it more manageable, since all of the data is stored at the same place.

10.2 Implementation

The SPI slave could have been implemented as a state machine. However because of the rather simplistic nature of the SPI communication protocol it was chosen not to implement it as a state machine. Instead it was implemented as a series of different events that trigger different responses in the SPI block. The first event is the S.CLK pin is having an rising edge. This event makes the the SPI slave read the input from the MOSI pin and prepare the next output for the MISO pin. The code that handles this event can be seen in figure 39.

```

if S_CLK_event = "01" then -- S_CLK rising edge
    if S_CLK_counter = 10 then -- Received all of address
        case input_shift_register(13 downto 11) is
            when "000" =>
                output_buffer <= "000" & CURRENT_POS_TILT; -- Swap output buffer with current position tilt
            when "001" =>
                output_buffer <= "000" & CURRENT_POS_PAN; -- Swap output buffer with current position tilt
            when others =>
                output_buffer <= (others => '0'); -- Set output buffer to default
        end case;
    end if;

    if S_CLK_counter /= -1 then -- If last entry, move to input buffer
        input_shift_register(S_CLK_counter) <= MOSI; -- Put data in shift register
        S_CLK_counter <= S_CLK_counter - 1; -- Decrement counter
    end if;
end if;

```

Figure 39: An S_CLK code example

The second event is when the SS pin has a rising edge. This signifies that the transmission is over and that all of the packet should be sent. This event triggers the SPI slave to update the data received on the correct output according to the received address for the data. An example of the code doing this can be seen in figure 40.

```

if SS_event = "01" then -- SS rising edge
    S_CLK_counter <= 13; -- Counter is reset

    case input_shift_register(13 downto 11) is -- Send data to address
        when "000" =>
            TARGET_POS_TILT_temp <= input_shift_register(10 downto 0);
        when "001" =>
            TARGET_POS_PAN_temp <= input_shift_register(10 downto 0);
        when "010" =>
            MAX_SPEED_TILT_temp <= input_shift_register(7 downto 0);
        when "011" =>

```

Figure 40: An SS code example

During the data transaction, two temporary registers are used, one for the input and one for the output. This is done to stabilize the data send so that it does not change mid transmission which would cause the data to become corrupt. There are 4 different kinds of settings that can be controlled via the SPI slave. The first setting is the target position used in the controller block, the second setting is the maximum speed for the motor, the third setting is the minimum speed of the motor and the last setting is the enabling of the motor. All of these 4 settings can be set for the pan and tilt part separately. This results in 8 different values needing an address resulting in an address length of 3 bits. Since the data length has to be constant, the data length was chosen according to the largest value needed to be sent. This was the target position, needing 11 data bits hence the data length was chosen to be 11 bits.

The output from the SPI slave to the SPI master was chosen to mimic the address and then send some data according to the address received under the transmission. This means that the output buffer has to be changed under the transmission. This would be a problem in most cases, however since the FPGA works much faster than the microcontroller this is not a problem. Since the only information the SPI master is interested in is the current position, the output buffer is only changed when this information is polled. It was made so that this information was always polled when a target position was sent, meaning that if only a poll on the current position is wanted, the target position is resent with the same value. In all other scenarios than the master sending the target position, the value returned to the master is the address the master sent and a data containing zeroes. The code handling this can also be seen in 39.

10.3 Tests

The SPI slave component has been tested in a Xilinx ISE testbench. This was done by simulating a transaction to the SPI slave. In figure 52 the test bench can be seen. SS can be seen going low signifying the start of the data transfer.

The S_CLK then starts to toggle and the data is shifted into the input register and as SS goes high the data from the input register is set to the placement the address is pointing to. In this case, a value to the target of the pan is being send. This has the address of 001 which can also be seen on the three first bits in the input register.




Figure 41: An Xilinx testbench simulation of the SPI

10.4 Discussion

One of the features you would get from creating the SPI slave component as a state machine is more robustness against errors. When an error happens in a state machine it would be unable to make the same amount of mess because it is locked to the state. However since no error should be happening and the fact that at no time while testing the connection was an error made, it is not really of any concern.

10.5 Conclusion

Both the testbench and physical test with the FPGA and microcontroller showed that this component functioned as it should.

11 Controller block

This section will describe in detail how the controller block receives a target position and from it generates a duty cycle and direction based on the feedback.

11.1 Description

As mentioned before, the goal of this block is to regulate the motors in order to reach a desired position accurately. This can be done by implementing a closed loop controller. This controller was implemented as two separate components, a comparator component that generates the error used by a **PI** controller component to generate a duty cycle related to the error. By making a customized controller and a separate comparator component, it was also easier to implement some features, like 'shortest path', which will be discussed later on. Apart from the reasoning given by the control theory section, another reason to use a PI controller is that it is

fairly easy to implement since the math required is simply gain and integration.

11.2 Mathematical techniques

This section will describe the theories used for the implementation.

11.2.1 Discrete Integration

Integration in discrete time is done by continuously sampling on a signal and relating the samples to each other. The technique used creates the square according to the previous error since this is arguably easier than other techniques to implement and that at a fast sampling rate the better precision gotten by using another technique is negligible. The method can be seen illustrated on figure 42. The equation for the method can be derived to the following: $Y = Y_0 + X_0t$. Where Y is the output, Y_0 is the previous output, X_0 is the previous sampled value and t is the time between each sample. This equation only makes sense in the context that it is seen as a process that is run each time you take a new sample.




Figure 42: Example of how integration is done.

11.2.2 Shortest path

Since the tilt part of the pan and tilt system is not restrained by wires, it is capable of turning more than 360 degrees. This means that the position can be related to that of a circle. However the control system sees this as straight line. This results in the control system being unable to go through the point of 0 if this is desired. The reason this happens is because a normal comparator is comparing the two values as if they were on a line instead of on a circle. So instead of making a normal comparator, a comparator that is capable of comparing values on a circle has to be made. In order to make such a comparator, an analysis of the differences between values on a line and values on a circle is made. The first difference is that on a line there is only one way to a point, however there are always two ways on a circle. A notable observation is that the directions on the circle will always be opposite of each other. The second difference, which is more like a similarity, is that one of the ways on the circle will always be equivalent to that of the line. The last thing to notice about the circle is that the distance is finite and since we know of the equivalence similarity, it can be derived that the distance of the other way has to be that of the equivalent part subtracted from circumference of the circle. It is also known that the direction of the other way

has to be the opposite of that of the equivalent part. In conclusion, it is possible to make a comparator working in circles by applying some basic mathematics to the result of a normal comparator. Comparing the length of the two ways will determine the shortest path. An illustration of the two ways you can go on the circle can be seen on figure 43.

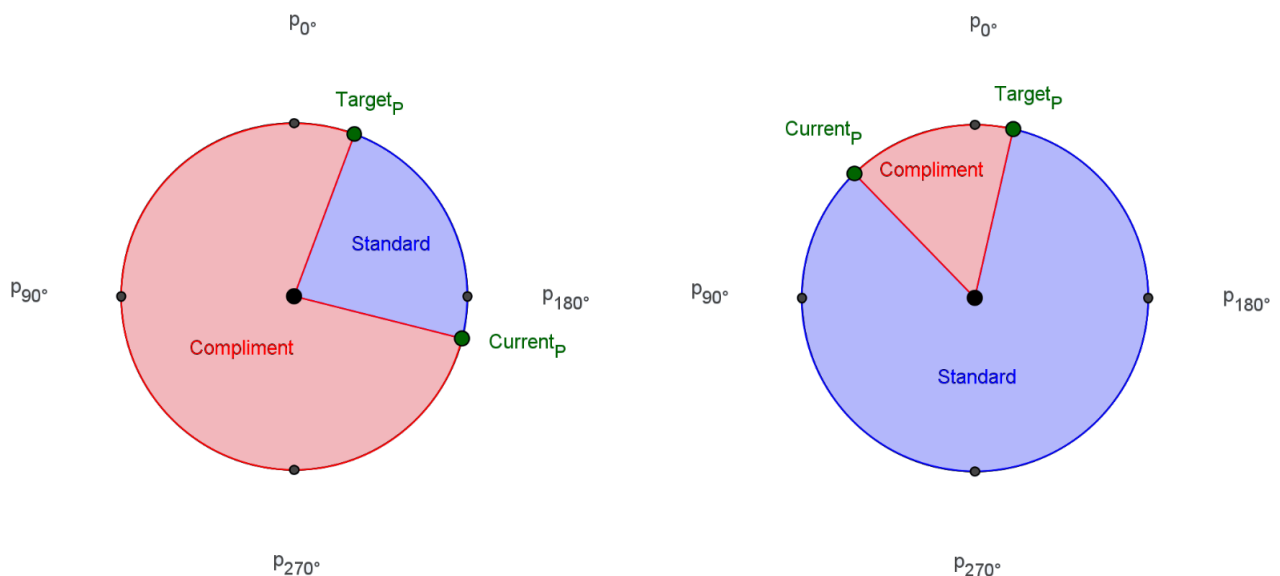


Figure 43: Illustration of how the shortest path is found.

11.3 Implementation

The implementation will be explained by looking at the comparator and the PI controller separately. The reason this can be done is that the only connection these two have is the error produced by the comparator.

11.3.1 Comparator

The comparator receives the target position from the SPI slave and the feedback from the sensor, then starts by doing a normal comparison. This is done in a way where no negative values are reached by looking at which of the two is the smaller and subtracting this from the larger one. The direction is then set according to this criteria as well. This comparison can be seen on picture 44.

```
compare_val <= INPUT - FEEDBACK when INPUT > FEEDBACK else FEEDBACK - INPUT; -- No negative values
complement_val <= conv_std_logic_vector(max_val, 11) - compare_val; -- The complement way on the circle
direction <= '1' when (INPUT > FEEDBACK) else '0'; -- Calculate direction
```

Figure 44: Example of code from the comparator.

The next step for the comparator is to determine which direction on the circle leads to the shortest path between current and target position, which is accomplished by using aforementioned mathematics. The length of the shortest path is then sent to the PI controller as the error and the direction is sent directly to the motor driver since the PI controller does not work in negative values, which will be discussed later. Since the pan is restrained by the cables, a safety feature is needed to make sure that if the pan receives a position that contradicts this restraint, it does not move there. This was achieved using a position constraint in the comparator module. Note that since this

is a base component that is shared by both the pan and tilt, it should be possible to turn off these constraints.

11.3.2 PI controller

The PI controller can be separated into calculating the P part, calculating the I part and then adding them together. Calculating the P part is quite simple since it is just taking the error and multiplying it by the constant set in the controller. However the I part requires that a sampling process is running. Inside this process a sampling is made and the arithmetic for the integration is done afterwards. The mathematics for the integration is simple enough to be implemented directly into code since t is a constant value. The sampling rate for the system was chosen to be 100kHz since this makes the error in the integration insignificant. The integration process can be seen on picture 45. Note that the division of t happens outside the process, since putting it in the process would make the value stay 0 because VHDL as a standard does not work in decimal and would just see values under zero, as zero. The summation is straightforward since it just requires P and I to be added. This concludes the controller part of the component, however there are still some practical restrictions that need to be added in order for the controller to work properly.

```
-- Integration --
integration_pro:
  process(clk_downscaled)
  begin
    if rising_edge(clk_downscaled) then
      if err > dead_band then -- Deadband check
        if saturation = '0' then -- Saturation check, no wind-up
          if (conv_integer(integrator(24 downto 0)) / (128*4)) < 45 then -- I limitation
            integrator <= integrator + prev_ERROR;
          end if;
        end if;
        if integrator < 0 then -- No negative integrator
          integrator <= (others => '0');
        end if;
      else
        integrator <= (others => '0'); -- Reset integrator since it is at
      end if;
      prev_ERROR <= ERROR_temp; -- Save previous error
      ERROR_temp <= ERROR;
    end if;
  end process;
```

Figure 45: Example of the code from the PI controller.

One of the restrictions needed is the output format. Since the motor driver component requires a duty cycle represented as a value between 0 - 255 (8 bits), an upper limit for the output value is set. If the controller creates a value above this limit, the output is simply set to the limit, ensuring no data loss from overflowing the 8 bits. To make the integrator more stable, anti wind-up was implemented. This was done by making the integrator halt the integration if the output is in saturation. This can be seen on the if statement checking for the saturation flag on picture 45. This ensures that the integrator is only running when the effect is needed.

Since the PI controller only works with positive values, this creates a problem where the integrator is only capable of increasing sum which essentially makes the system unstable. The best solution would always be to make the integrator integrate negatively. This would however require us to work in negative values. The integrator was implemented where it would reset upon having an error of 0. This makes the controller less naturally close to the target position, however since we are tuning the PI controller not to overshoot, this less natural behaviour is negligible. Doing it this way would also take away some of the potential overshoot. Upon testing the controller, it was noticed that integrator was reacting quite slowly, possibly because of the wind-up protection, which made it

stop before reaching the destination and after a short time of standstill, the integrator would kick in, making the motor move the last bit. This could be solved simply by increasing the constant for the integrator part, however this would make it overshoot which was not desired. To solve this issue, a limit on the integrator part was implemented so that you would be able to increase the response of the integrator without the integrator making the system overshoot. The limit basically limited the speed the integrator part is contributing with.

11.4 Tests

The components have been tested separately because of the same reason stated before. The comparator tests were mainly done through simulation via a testbench in the Xilinx IDE and the PI controller test were done on the pan and tilt system.

11.4.1 Comparator

The comparator was tested for different values which resulted in table 46, showing the tested values, the expected result and the actual result.

Input	Feedback	Container	Expected value	Simulated value	Expected Direction	Simulated direction
100	500	off	400	400	0	0
800	500	off	300	300	1	1
700	100	off	480	480	0	0
100	500	on	140	140	0	0
800	500	on	210	210	1	1
700	100	on	600	600	1	1

Figure 46: Comparator test table.

11.4.2 PI controller

The PI controller were tested manually on the pan and tilt system. This test was done after the other components were tested working. A video of the test can be seen on Appendix B.

11.5 Discussion

Some of the practical solutions in the PI controller have the consequence that the PI controller does not entirely follow control theory. Some of these problems would be not be there if the controller worked with negative values and since later in the development stage it was discovered that implementing the controller as a filter would have been optimal solution, these problems would not exist.

11.6 Conclusion

The components work as expected and the controller regulates as desired.

12 Sensor block

This section describes in detail how the sensor block handles the the output of the Hall effect sensors and sends it to the comparator block.

12.1 Description

The responsibility of the sensor block is to handle the output from the Hall effect sensors and determine the position of both the pan part and the tilt part of the system. The controller block rely on this information and thus this block form the basis of the control method chosen.

12.2 Theory

Hall effect sensors are sensors that change their voltage output based on proximity of a magnetic field. If combined with the right circuitry it can be used a switch, giving a digital output. When placed around the rotating axis of a motor with the right relative distance, the Hall effect sensors can be used determine the position, direction, rotational velocity and the position of the axis. There is a disc on the axis of the motor. It is divided into 6 sections, alternating between two levels of magnetization. Around the magnet, two Hall effect sensors are placed at an angle of slightly less than 90 degrees(See figure 47). As the motor axis rotate, the magnetized parts move past the sensors and the magnetic field, they are subjected to, changes. Because the Hall effect sensor is a little “off” the output will have an overlap, see figure 48. This is what is used to determine the direction.

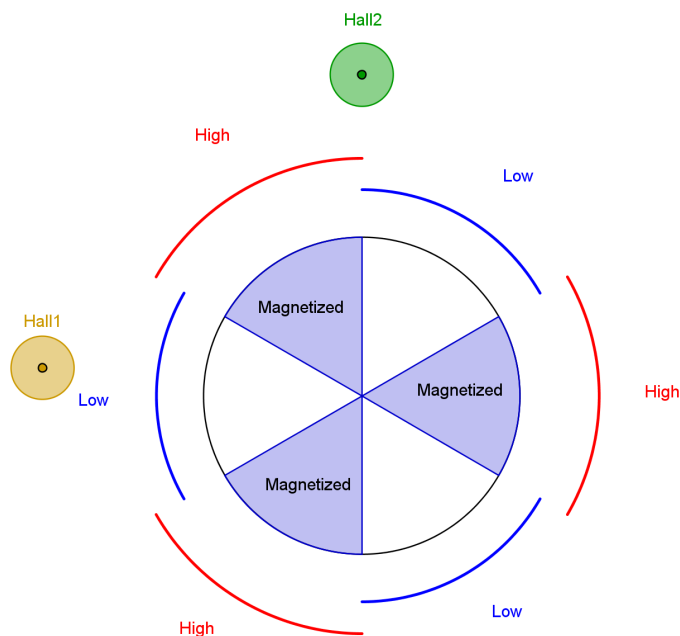


Figure 47: This is an example of how the Hall effect sensors is placed around the disk.




Figure 48: As the axis rotates, the output changes

12.3 Implementation

The Hall counter is the sensor block. Every time the Hall effect sensors gives an output the Hall counter is updated. There are 1080 “tachs” or changes in the output of the Hall effect sensor, for one revolution. To be able to move both clockwise and counterclockwise, the starting position is set to half of the total: 540. This is done to avoid working with negative values, as that can contribute to unnecessary difficulty in the implementation.

It starts out with sampling the value of the sensors. The value of the sensor is read and shifted into a vector of two bits as seen in figure 49

```
if rising_edge(CLK) then -- Sample and shift
    debounce_hall1 <= debounce_hall1(0) & HALL1;
    debounce_hall2 <= debounce_hall2(0) & HALL2;
end if;
```

Figure 49: An example of how the value is sampled and shifted

If no change has been registered, the position doesn't change and a new sample is taken. As seen in figure 50. If a change is detected then a series of “if statements” determines whether the change in position is positive or negative.




Figure 50: Flowchart of how the position changes in relation to different Hall effect sensor outputs

A small example of the code involved can be seen in figure 51

```
if falling_edge(CLK) then -- Calculate pos
    if debounce_hall1 = "01" then
        if HALL2 = '1' then
            pos := pos + 1;
        else
            pos := pos - 1;
        end if;
    end if;
```

Figure 51: Small bit of code that shows how the position is calculated based on the Hall effect sensor output

12.4 Test

To make sure the code behaves in the intended way a test bench has been made. It simulates the code's response to changing output from the Hall effect sensors. The first two bars show the Hall effect sensors and their changing values. The next bar is the clock and under that one the position.

As can be seen in figure 52, the position is steadily decreasing in value. If the position reaches the max value (108010 or 00000100001110002) the position is set to 0. If the position reaches -1 it is set to the max value minus one.




Figure 52: Testbench of how the position changes based on Hall effect sensor outputs

12.5 Discussion

As there is only two Hall effect sensors and the max rate of change in their output is a lot less than the possible sampling rate, it is fairly easy to calculate the change in position of the motor. However having only two sensors and only 6 different magnetization levels limits the accuracy of the measurements. If the axis of the motor moves in the area between a rising or falling edge but never reaching one on either sensors the change in position cannot be measured. More sensors would give a higher precision but also require more calculations in order to determine whether the change is positive or negative. One more error that can be experienced is if the axis end up in a position where one of the sensors is “stuck” on a rising or falling edge. If that happens the position will constantly flip between the two values. Lastly as magnets are used they sometimes “stick” to a certain position. All this combined with the slack in the belt transferring the rotation from the motor to either the pan or tilt and the small distance before the teeth of the belt and gear meshes, this limits the overall precision of the system

12.6 Conclusion

The component works as expected and sends the current position of the system to the comparator block.

13 Motor driver block

This section will describe in detail how the motor driver block receives a duty cycle and a direction, then uses these to generate and output a **PWM** signal according to the direction and speed set by the controller block.

13.1 Description

An easy way to visualize the inner workings is to divide the Motor Driver into two separate blocks. The first block's task is to create a **PWM** signal which is equivalent to the duty cycle given by the controller. The second block's task is to set the output of the **PWM** signal to the right pin on the H-Bridge in order to control the direction of the motor.

13.2 Theory

This section will describe the theories used for the implementation.

13.2.1 PWM generation

The **PWM** signal is created based on a looping counter and a threshold. Imagine a counter resetting whenever it reaches a certain value, creating a loop. If the value is set to the period of the wanted **PWM** signal and the **PWM** signal output is dependent on the counter value relative to a threshold, then you will be able to create a looping output that is dependant on the threshold. Figure 53 shows an illustration of the behavior of **PWM** signal based on the looping counter and threshold.




Figure 53: An example of a H-bridge

13.2.2 H-bridge

An H-bridge is a collection of 4 transistors, designed to control a DC motor. In figure 54 the concept is illustrated. When the red switches are on, the motor runs in one direction and vica versa for the blue switches. There are two H-bridges in the system, one for the pan motor and one for the tilt motor. For further information about H-bridges, consult the datasheet of the H-bridge in appendix B.




Figure 54: An example of a H-bridge

13.3 Implementation

The implementation was done with the separation of the block described before. A **PWM** component was created which has the sole purpose of creating a **PWM** signal based on a duty cycle. The other part has the purpose of directing the **PWM** signal to the right output pins based on the direction wanted.

13.3.1 PWM part

The PWM block was implemented so that whenever the counter value is below the threshold the PWM signal output is high and when the counter value is above the threshold the PWM signal output is low. The section of the PWM code that makes the counter can be seen on picture # Counting example #.

13.3.2 Direction part

This block is simply implemented as a switch that toggles on the wanted direction. As an example if the wanted direction is set to 0 then the PWM is driven to the relevant motor control output and the other is kept at zero.

13.4 Tests

To verify that this motor controller is working as intended, a test bench was created and various inputs was tested. The enable was set and a duty cycle of 8 bits was input, which corresponds to a certain duty cycle. A clear spike in PWM output could be seen, of the same length every time PWM is outputting. To further verify the test, the edges of the PWM signal was measured and used to find the exact duty cycle. Figure 55 showcases the tests of PWM.

Duty Cycle	PWM duration	PWM period	Calculated Duty Cycle	Measured Duty Cycle
"0000 0000"	1 ns	1281 ns	0.00 %	~0.00 %
"0010 1111"	236 ns	1281 ns	18.43 %	18.42 %
"1010 0000"	801 ns	1281 ns	62.75 %	62.53 %
"1111 1111"	1276 ns	1281 ns	100 %	99.61 %

Figure 55: PWM Tests

13.5 Discussion

The tests show a deviation from the desired duty cycle of a small percentage. However these deviations have been deemed insignificant. The errors are created by the method that is used to calculate the PWM from the duty cycle. When converting the bit string, it is multiplied by the PWM period and divided by 256. This yields an integer, that when combined with the counter, yields a PWM very close to the desired duty cycle. Beyond that margin of error, a single clock cycle is lost due to resetting the counter, which corresponds to 1ns.

13.6 Conclusion

The motor driver has been sufficiently explained and tested, and has become a solid module to create a PWM signal for the P & T system. The PWM calculation can be optimized very slightly by improving the conversion from duty cycle to PWM, and optimizing the counter, as to never lose any clock cycles.

14 Driver conclusion

Even though the driver runs quite smooth some changes could have been made in order to make it better. One of these changes would be to make a controller that actually follows the control theory for a PI controller a little better. The reason why this was not done from the start, was because of a lack of understanding how a controller could be implemented. As an example it was discovered rather late in the development stage that the optimal way

of implementing the PI controller would have been by making a filter. This was however discovered too late to change. A filter was created, but an unknown error shut down the idea of implementing it as a filter since the time did not allow it to be completed. However if more time was given implementing the controller as a filter would be one of the upgrades that would make the product easier to tune because the filter follows control theory closer than the implemented controller. Another thing would be to implement the controller as a PID controller instead since this would make a smoother controller. The control theory section also derives that the optimal controller would be the PID. The easiest way to implement the D part would be to make a speed sensor, this was however discovered late in the development stage after a failed attempt at implementing the D part through standard derivative mathematics. An error was also found in the link between the controller block and the motor driver block since the controller is actually based on controlling a voltage rather than the duty cycle. This error can be fixed by making a conversion from voltage to duty cycle in the controller block.

15 Conclusion

A controller based on a mathematical system model has been implemented on the FPGA. This controller works, but does not completely fulfill the requirements for overshoot and settling time. This was found to be because of real-world constraints, such as maximum and minimum voltage for the motor and unlinear behavior for the P&T systems. It does, however, fulfill the steady-state requirement, reaching the desired position with great accuracy.

The controller uses PWM to dictate the motor speed and an H-bridge to control motor direction, which has been tested to work perfectly.

This controller can communicate with a developed application through SPI. This communication supports a bitrate of 107k bits/sec, which is more than enough for the intended purpose.

The application creates an interface for a user to manipulate constraints and the position of the system. However, an important algorithm used for the conversion between stage-dimensions and P&T angle position has not been implemented in the application, due to time constraints, even though it has been derived. It is therefore plausible that the application can fulfill the height and stage-dimension requirements, once this feature is implemented.

As for the secondary requirements, only the application was implemented, as working on the primary requirements was deemed more important.

Because of time constraints and minor shortcomings, the system as a whole was not able to be tested extensively. However, the individual parts were all tested or mathematically proven to work.

16 Appendix

A Kernel Instruction List

Table 1: Instruction list

Function	Opcode (char)	Parameters	Notes
Connection_Check	o	0	Returns a message to the console
Enable_FPGA	f	1	Enable = 1, Disable = 0, by default FPGA is enabled
Goto_Coordinate	g	2	Updates current target in MCU, and sends a new input to the FPGA
Set_Height	h	1	Sends height constraint to MCU
Set_Width	w	1	Sends width constraint to MCU
Set_Depth	d	1	Sends depth constraint to MCU
Set_Scene	s	1	Select a predefined scene, and update MCU to those constraints
Set_Min_Vel_Pan	p	1	Sends minimum speed setting for pan to FPGA
Set_Max_Vel_Pan	P	1	Sends maximum speed setting for pan to FPGA
Set_Min_Vel_Tilt	t	1	Sends minimum speed setting for tilt to FPGA
Set_Max_Vel_Tilt	T	1	Sends maximum speed setting for tilt to FPGA
Run_Lightshow	l	1	Starts lightshow corresponding to parameter value
Stop_Lightshow	c	0	Stops lightshow
Enable_Joystick	j	1	Starts joystick user input

B Content on CD

1. Microcontroller Code
2. FPGA driver Code
3. FPGA Driver Xilinx Project
4. Cortex M4 Datasheet
5. Pan & Tilt Test Video
6. H-Bridge Data sheet
7. Project Description
8. Copy of Report